



北京航空航天大学
COLLEGE OF SOFTWARE 软件学院
BEIHANG UNIVERSITY

Chapter 7

GPU Parallel Programming

CUDA

软件学院 邵兵

2022年4月19日

Contents

- 1. Programming Model**
- 2. CUDA Language**
- 3. Example Code Study**
- 4. CPU & GPU Synchronization**
- 5. Multi-GPU**
- 6. Dynamic Parallelism**



What is CUDA?

CUDA: Compute Unified Device Architecture

- **CUDA is a **compiler** and **toolkit** for programming NVIDIA GPUs**
- **Enable heterogeneous computing and horsepower of GPUs**
- **CUDA API extends the C/C++ programming language**
- **Express SIMD parallelism**
- **Give a high level abstraction from hardware**



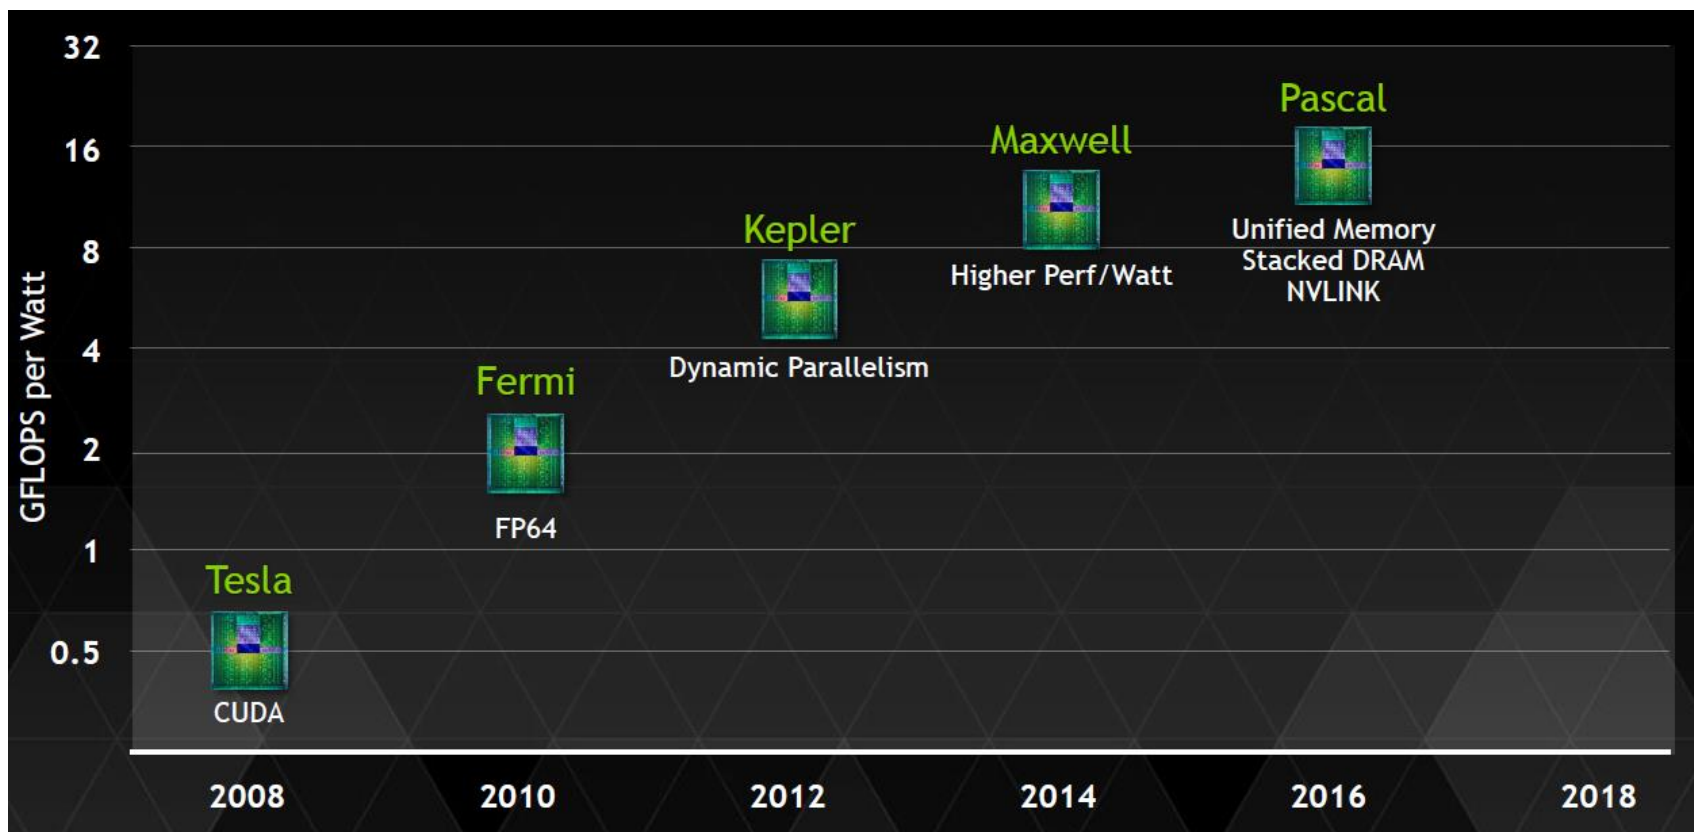
CUDA Version

CUDA SDK Version	Compute Capability	Architecture
6.5	1.X	Tesla
7.5	2.0-5.x	Fermi, Kepler, Maxwell
8.0	2.0-6.x	Fermi, Kepler, Maxwell, Pascal
9.0	3.0-7.x	Kepler, Maxwell, Pascal, Volta
10.2	6.1x	... Turing

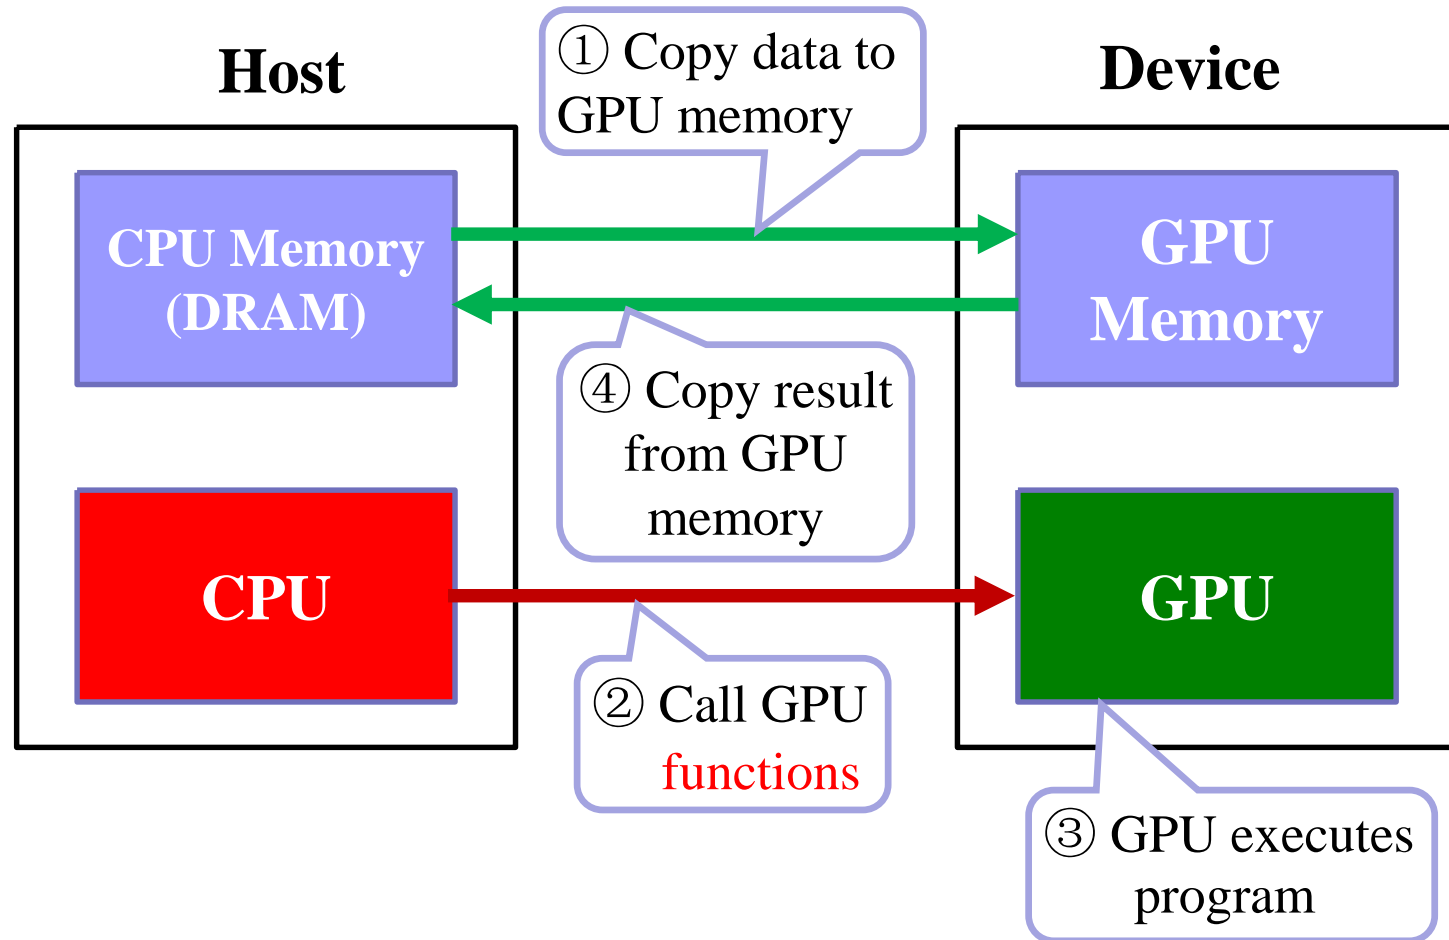
<https://www.cnblogs.com/timlly/p/11471507.html>



北京航空航天大学
COLLEGE OF SOFTWARE
BEIHANG UNIVERSITY 软件学院



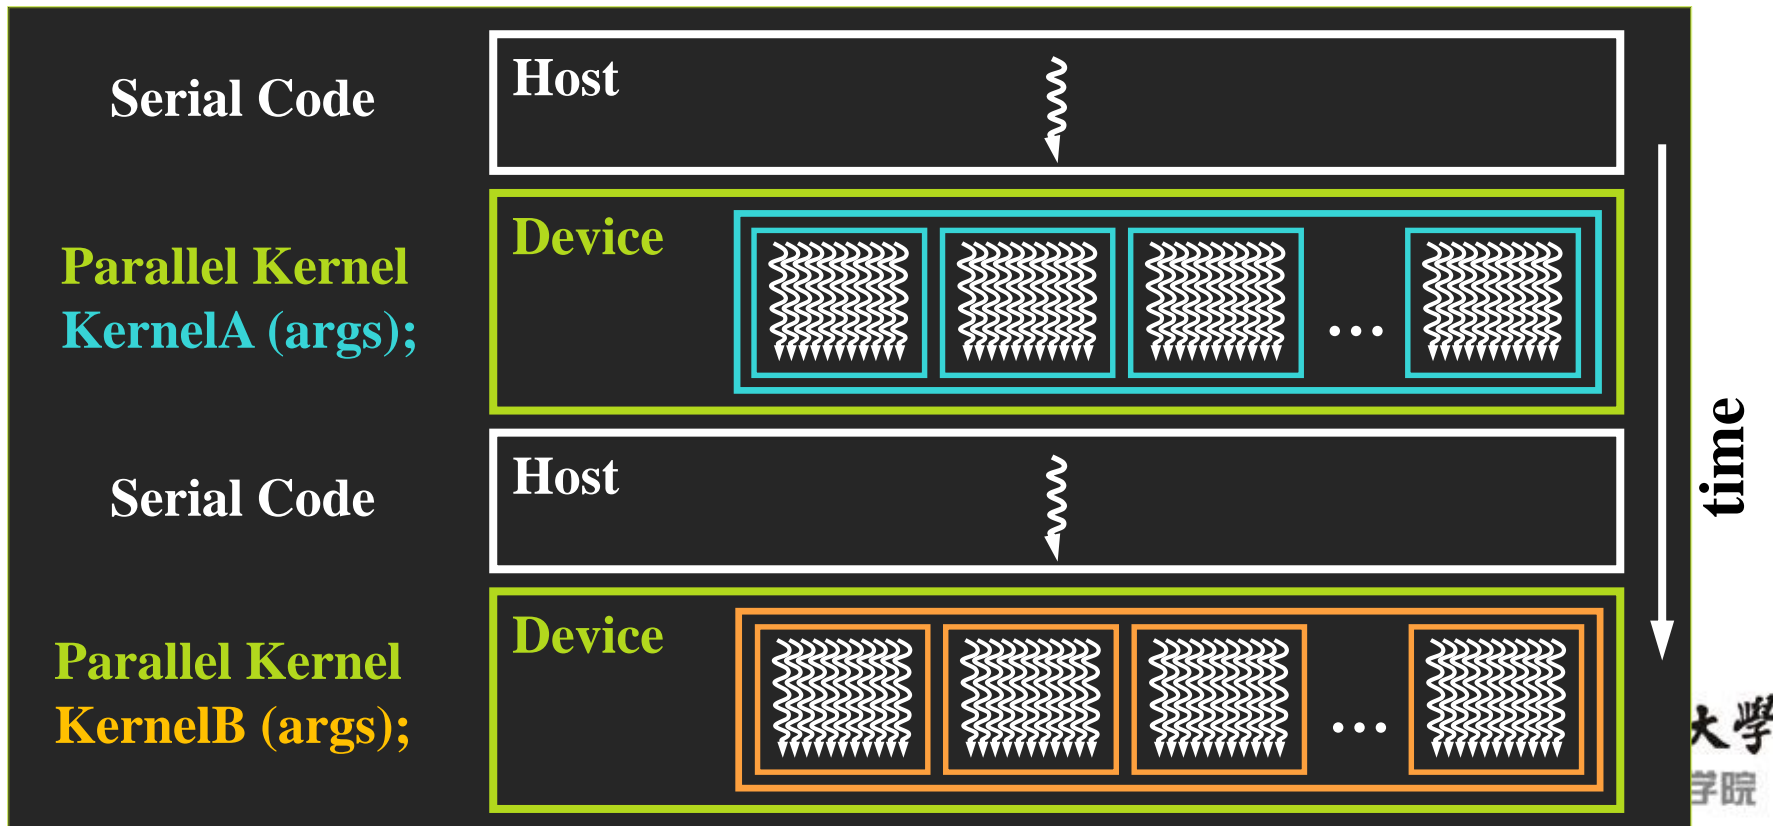
CUDA Program Flow



CUDA Programming Model

CUDA = serial program with parallel kernels, all in C

- Serial C code executes in a host thread (i.e. CPU thread)
- Parallel kernel C code executes in many device threads across multiple processing elements (i.e. GPU threads)



CUDA Program Framework

GPU code

```
#include <cuda_runtime.h>
```

```
__global__ void my_kernel (...) {  
    ...  
}
```

CPU code
(serial or
parallel if
pthread/
OpenMP/
MPI is
used.)

```
int main() {  
    ...  
    cudaMalloc (...)  
    cudaMemcpy (...)  
    ...  
    my_kernel<<<nblock, blocksize>>> (...)  
    ...  
    cudaMemcpy (...)  
    ...  
}
```



Kernel = Many Concurrent Threads

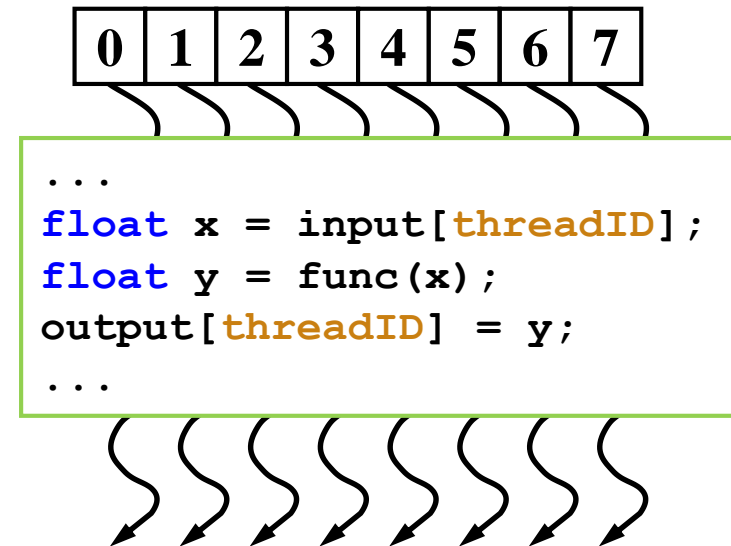
- One kernel is executed at a time on the device
- Many thread execute each kernel
 - Each thread executes the same code
 - ... on the different data based on its threadID

- **CUDA thread might be**

- **Physical threads**

- As on NVIDIA GPUs
 - GPU thread creation and context switching are essentially free

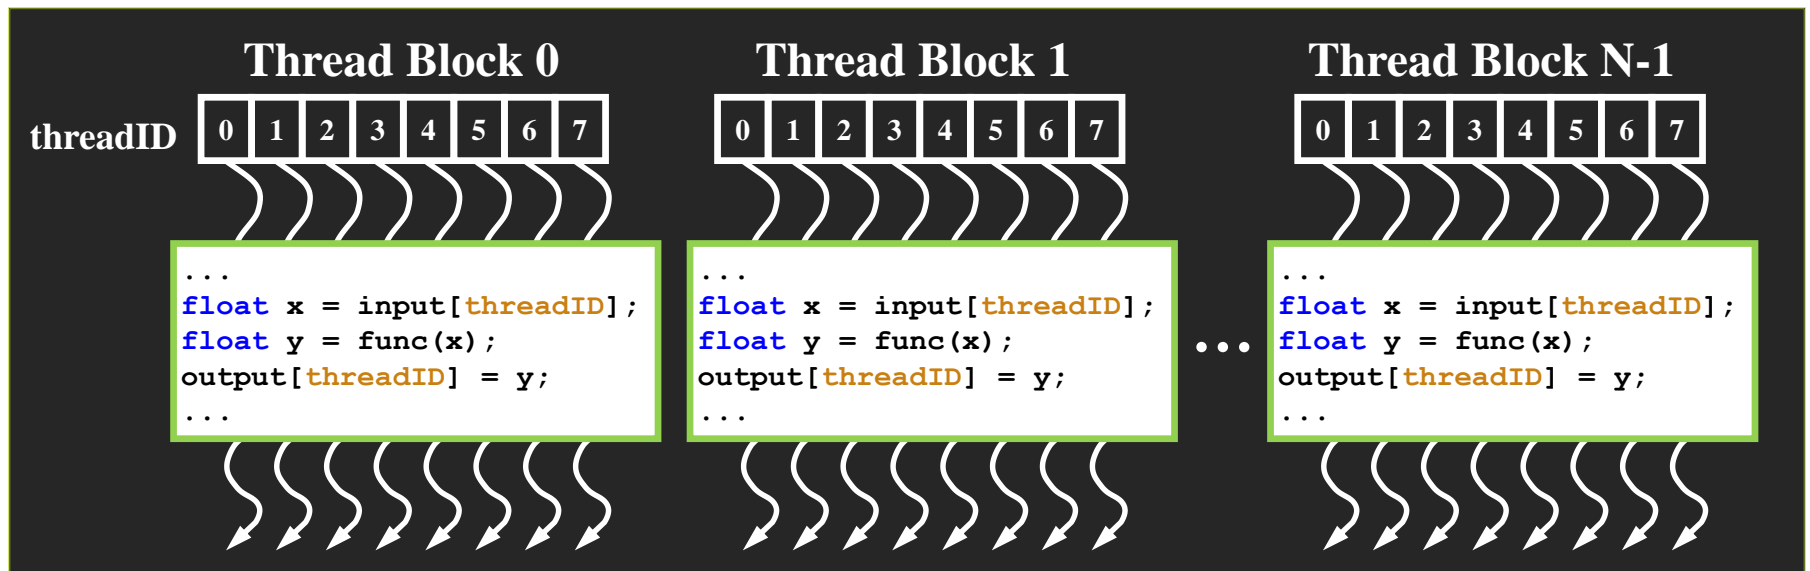
- Or virtual threads



E.g. 1 CPU core might execute multiple CUDA threads

Hierarchy of Concurrent Threads

- Threads are grouped into thread blocks
 - Kernel = grid of thread blocks



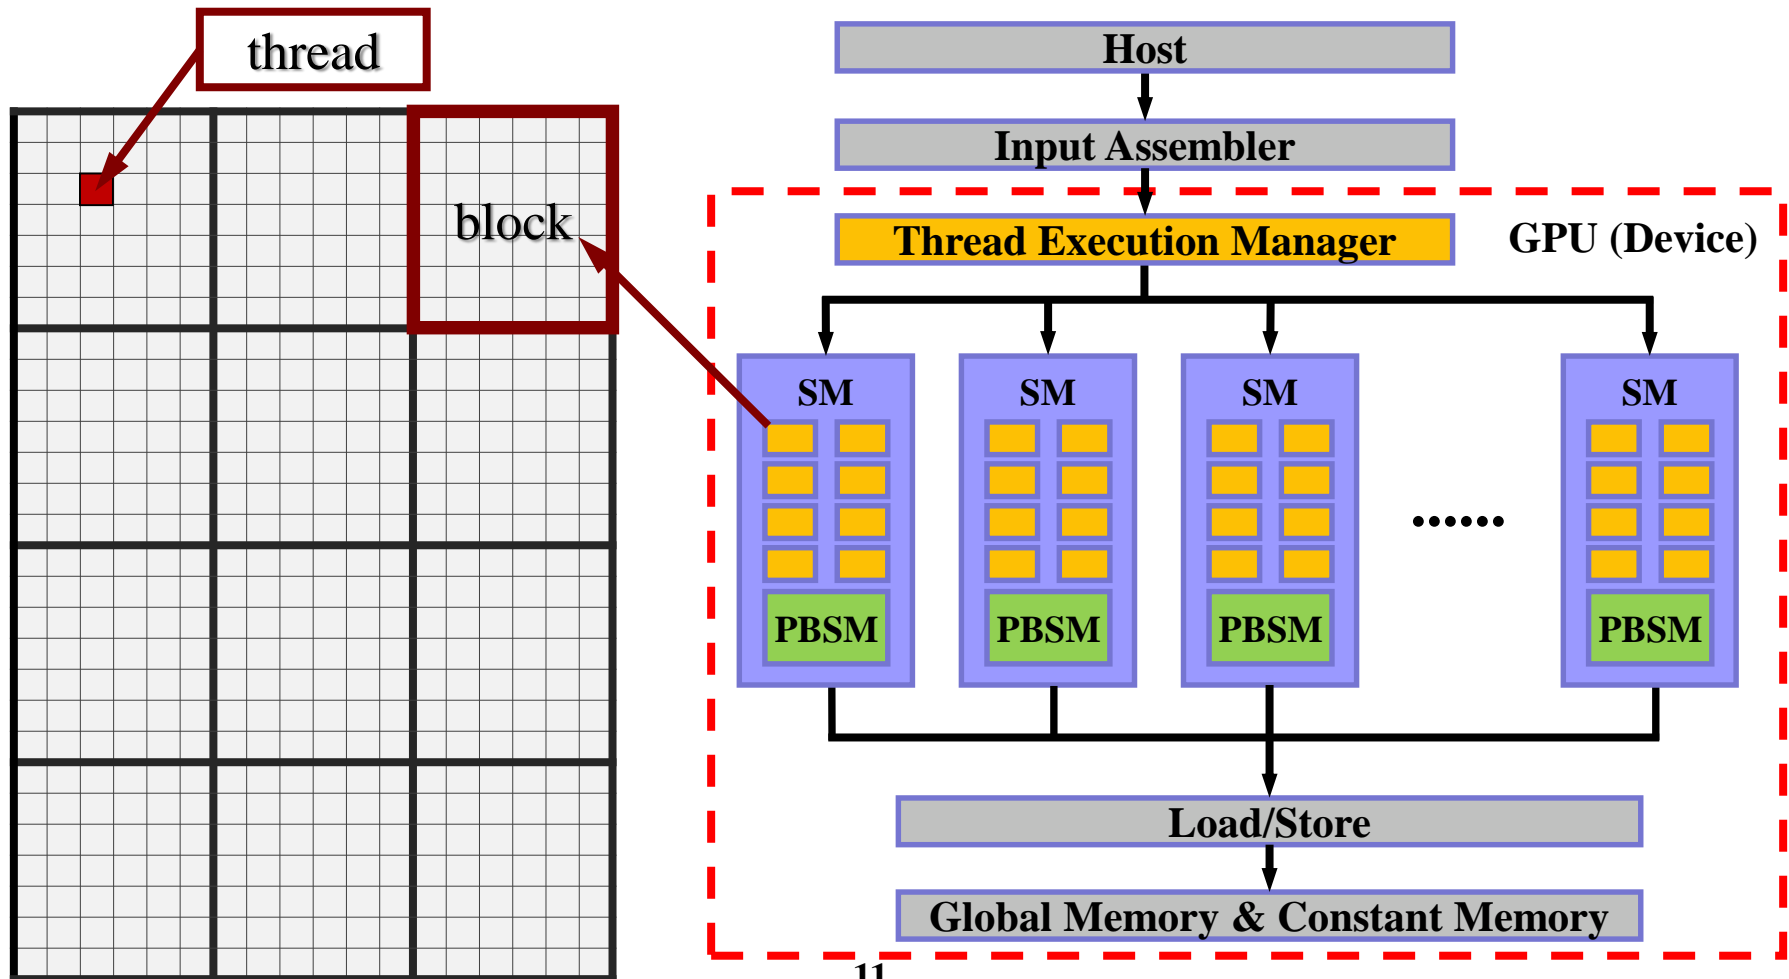
- By definition, threads in the same block may synchronized with barriers, but not between blocks.

```
scratch[threadID] = begin[threadID];  
__syncthreads();  
int left = scratch[threadID - 1]
```



Software Mapping

- **Software:** grid → blocks → threads
- **Hardware:** GPU(device) → SM(Stream Multiprocessor) → core



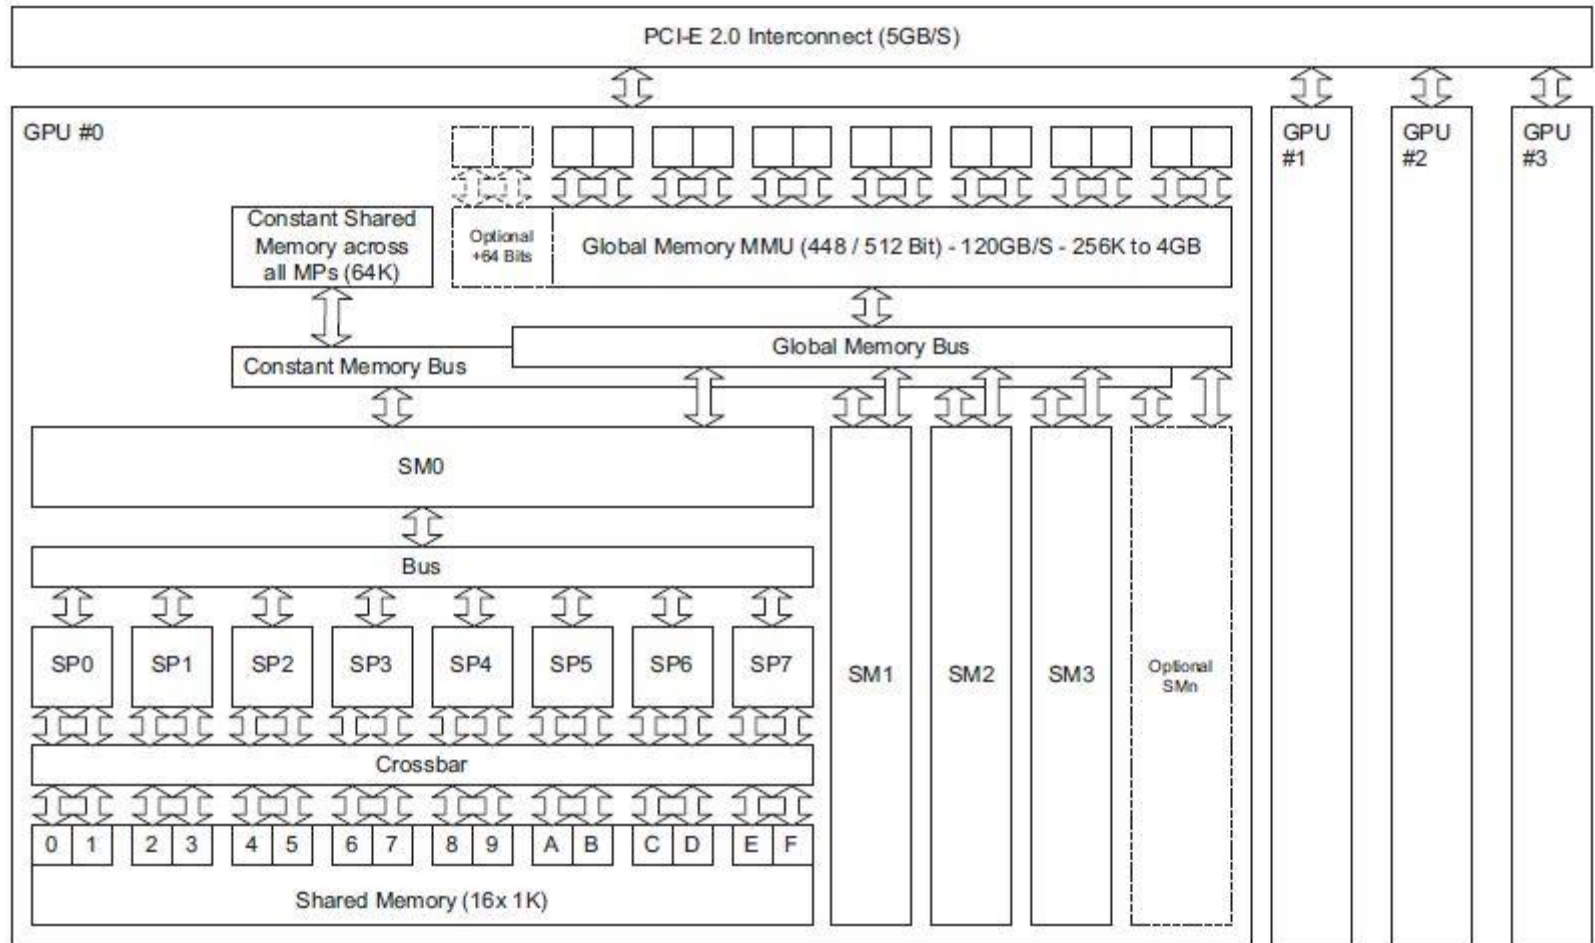
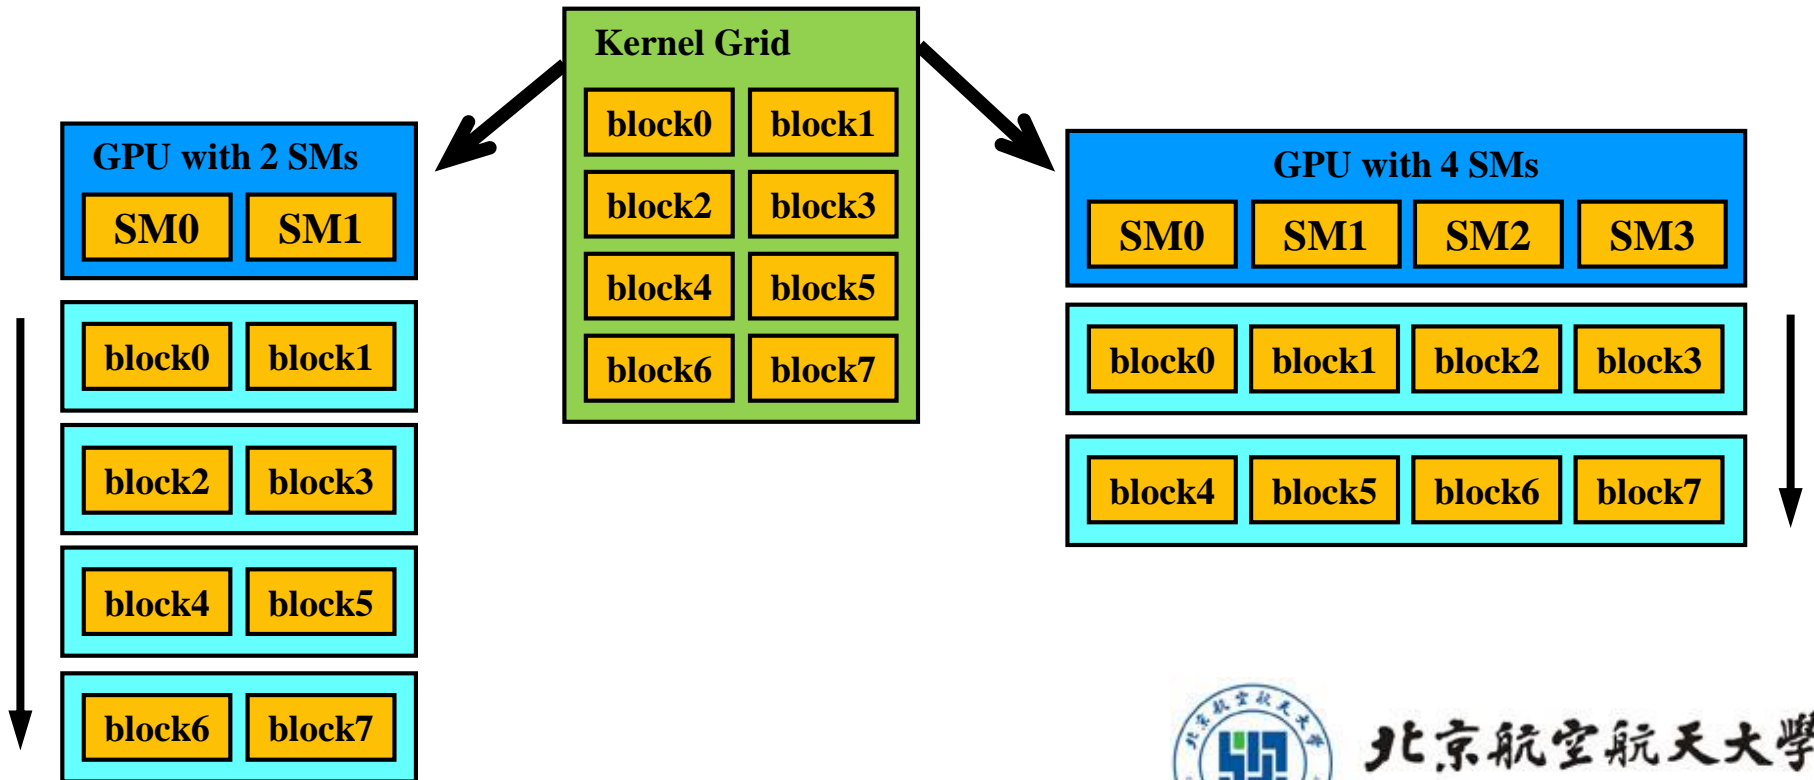


FIGURE 3.5



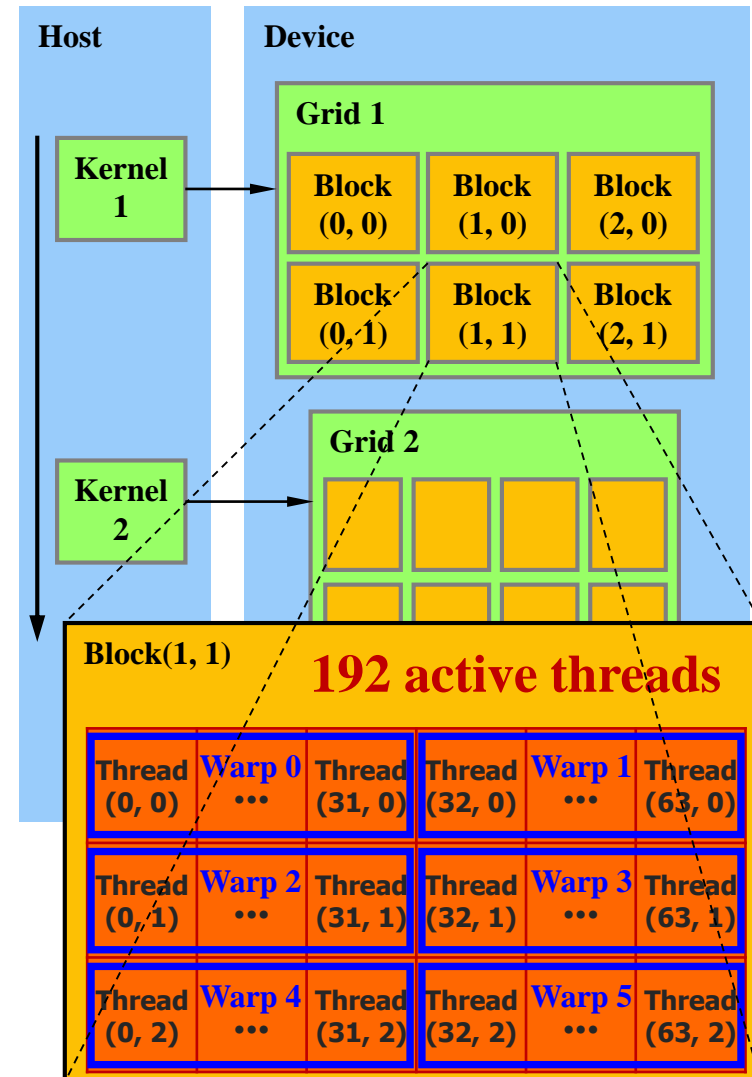
Block Level Scheduling

- **Blocks are independent to each other to give scalability**
 - A kernel scales across any number of parallel cores by scheduling blocks to SMs

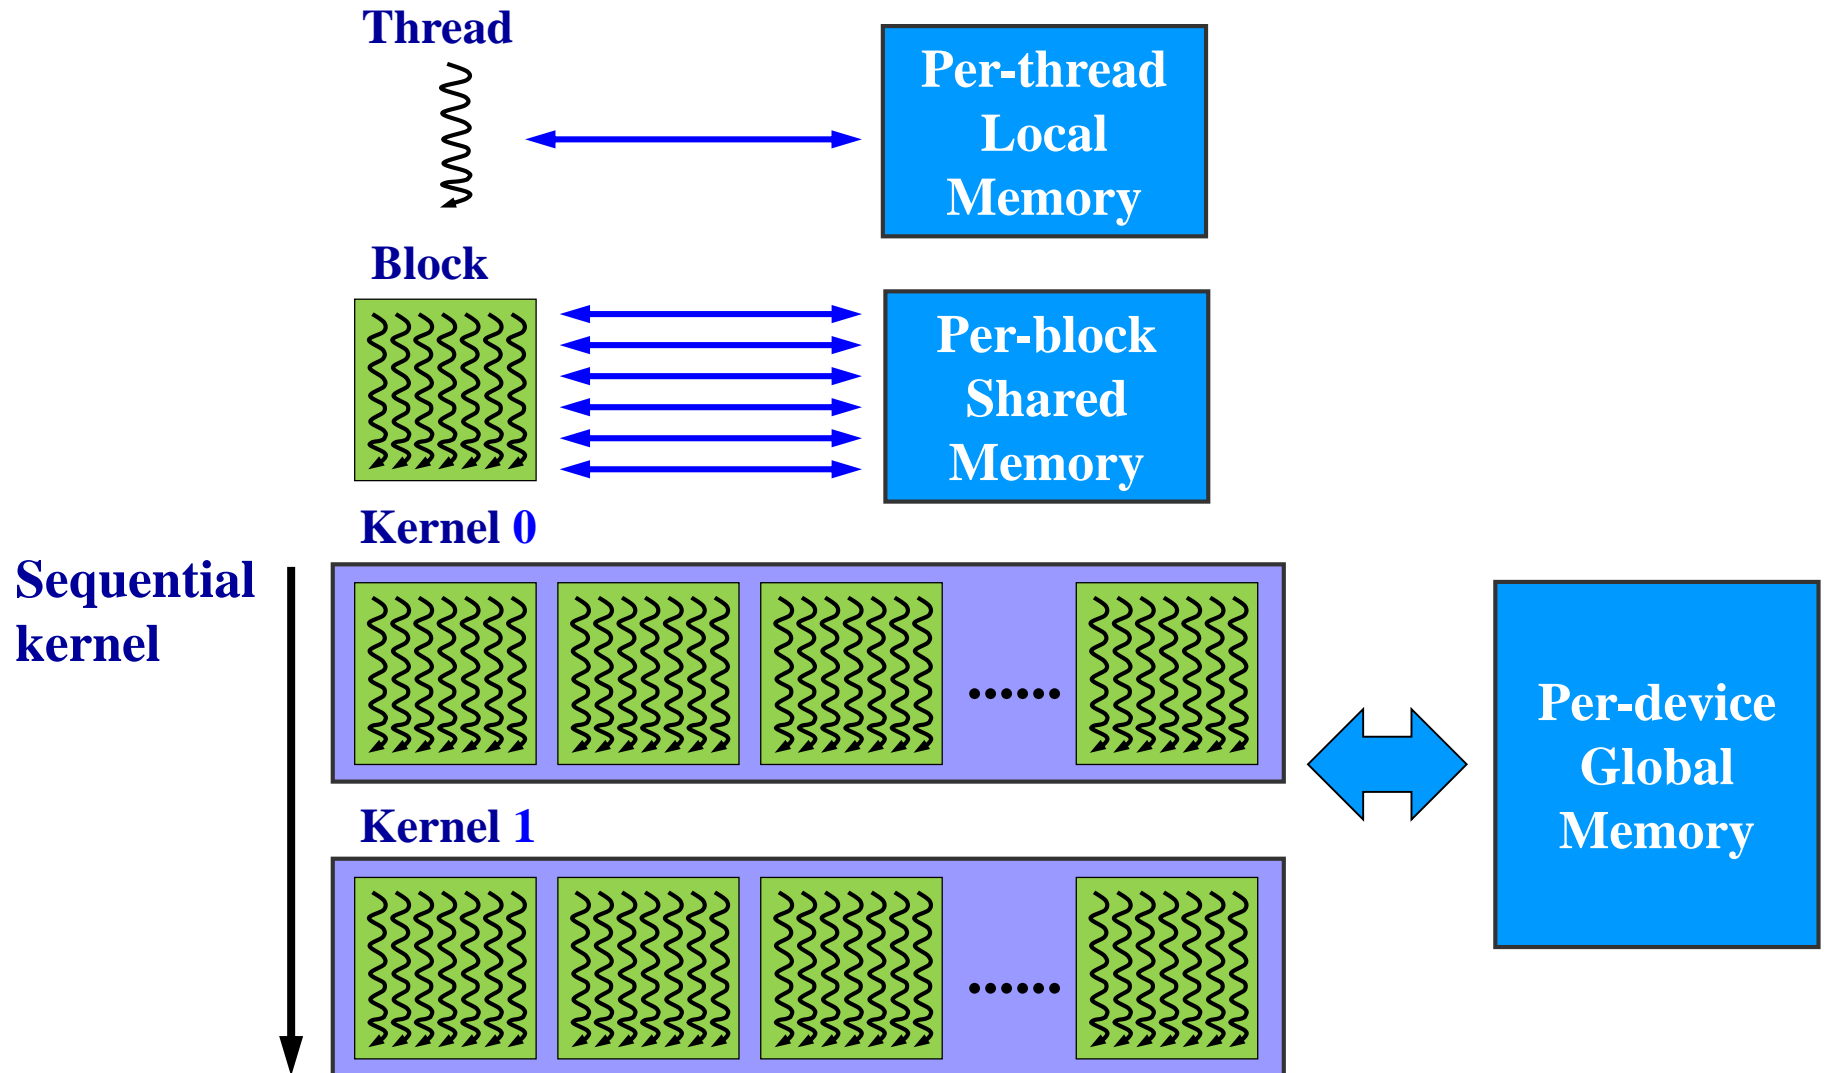


Thread Level Scheduling - Warp

- Inside the SM, threads are launched in groups of 32, called **warps**
 - Warps share the control part (**warp scheduler**)
 - Threads in a warp will be executing the same instruction (SIMD)
- In other words ...
 - Threads in a warp execute **physically in parallel**
 - Warps and blocks execute **logically in parallel**

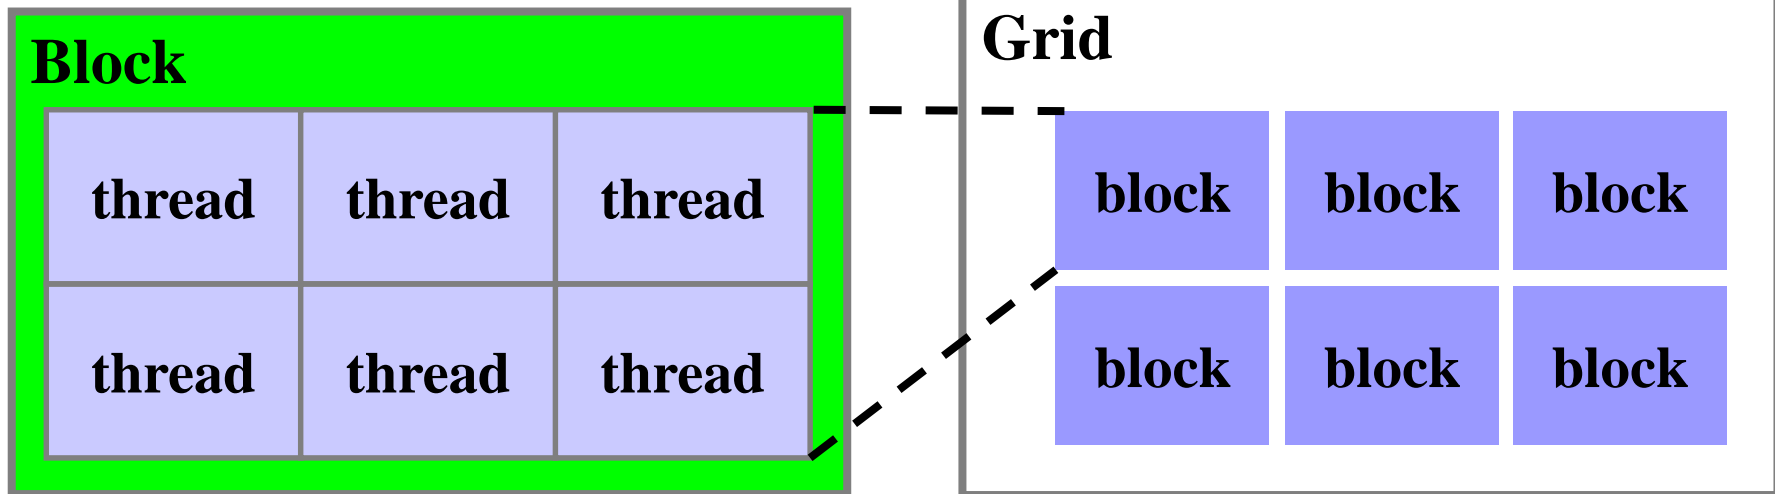


Memory Hierarchy



CUDA Programming Terminology

- Host : CPU
- Device : GPU
- Kernel : functions executed on GPU
- Thread : the basic execution unit
- Block : a group of threads
- Grid : a group of blocks



Contents

- 1. Programming Model**
- 2. CUDA Language**
- 3. Example Code Study**
- 4. CPU & GPU Synchronization**
- 5. Multi-GPU**
- 6. Dynamic Parallelism**



CUDA Language



Philosophy: provide minimal set of extensions necessary

■ Kernel launch

```
kernelFunc<<<nB, nT, nS, Sid>>> (...); // nS and Sid are optional
```

- nB: number of blocks per grid (grid size)
- nT: number of threads per block (block size)
- nS: shared memory size (in bytes)
- Sid: stream ID, default is 0

■ Build-in device variables

```
threadIdx; blockIdx; blockDim; gridDim
```

■ Intrinsic functions that expose operations in kernel code

```
__syncthreads();
```

■ Declaration specifier to indicate where things live

```
__global__ void KernelFunc (...); // kernel function, run on device  
__device__ void GlobalVar;         // variable in device memory  
__shared__ void SharedVar; // variable in per-block shared memory
```

Thread and Block IDs

- Build-in device variables

`threadIdx`; `blockIdx`; `blockDim`; `gridDim`

- The index of threads and blocks can be denoted by a 3 dimensional struct

➤ `dim3` defined in `vector_types.h`

```
struct dim3 {x; y; z;};
```

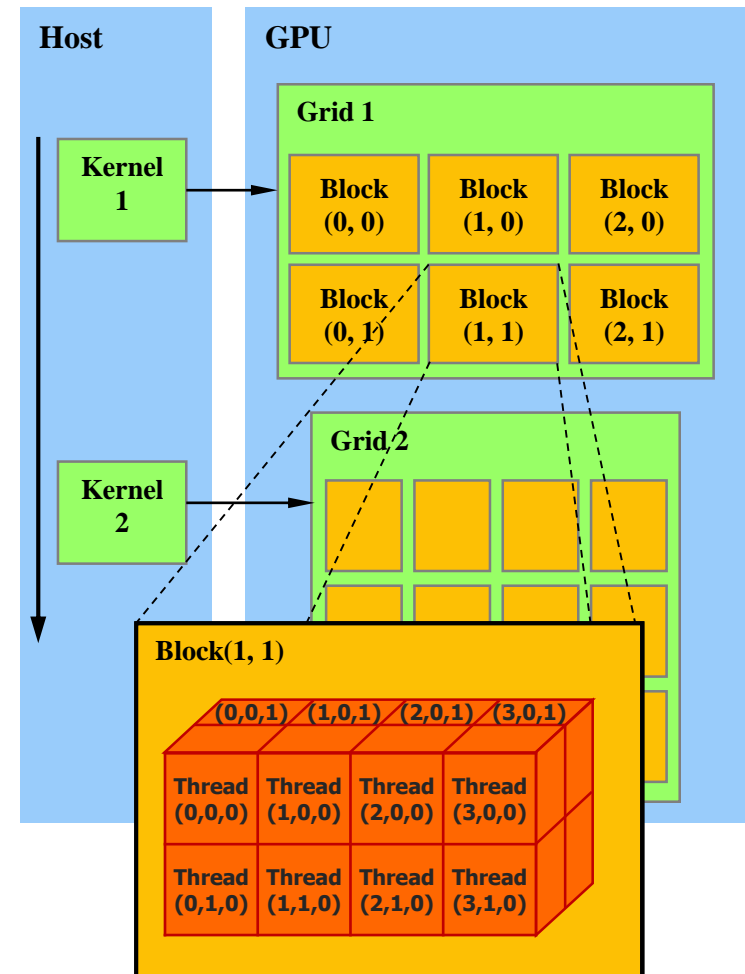
- Example

➤ `dim3 grid(2, 3);`

➤ `dim3 blk(3, 5);`

➤ `my_kernel<<<grid, blk>>>();`

- Each thread can be uniquely identified by a tuple of index (x, y) or (x, y, z)



Function Qualifiers

Function qualifiers	limitations
<code>__device__</code> function	Executed on the device Callable from the device only
<code>__global__</code> function	Executed on the device Callable from the host only (must have void return type!)
<code>__host__</code> function	Executed on the host Callable from the host only
Functions without qualifiers	Compiled for the host only
<code>__host__ __device__</code> function	Compiled for both the host and the device



Variable Type Qualifiers



Variable qualifiers	limitations
<code>__device__</code> var	<ul style="list-style-type: none">● Resides in device's global memory space
<code>__constant__</code> var	<ul style="list-style-type: none">● Has the lifetime of an application● Is accessible from all the threads within the grid and from the host through the runtime library● Resides in device's constant memory space
<code>__shared__</code> var	<ul style="list-style-type: none">● Resides in the shared memory space of a thread block● Has the lifetime of the block● Is only accessible from all the threads within the block



Device Memory Operations

■ Three functions

`cudaMalloc()`, `cudaFree()`, `cudaMemcpy()`

➤ Similar to the C's `malloc()`, `free()`, `memcpy()`

1. `cudaMalloc(void *devPtr, size_t size)`

➤ `devPtr`: return the address of the allocated device memory

➤ `size`: the allocated memory size (bytes)

2. `cudaFree(void *devPtr)`

3. `cudaMemcpy(void *dst, const void *src, size_t count, enum cudaMemcpyKind kind)`

➤ `count`: size in bytes to copy



cudaMemcpyKind

- one of the following four values

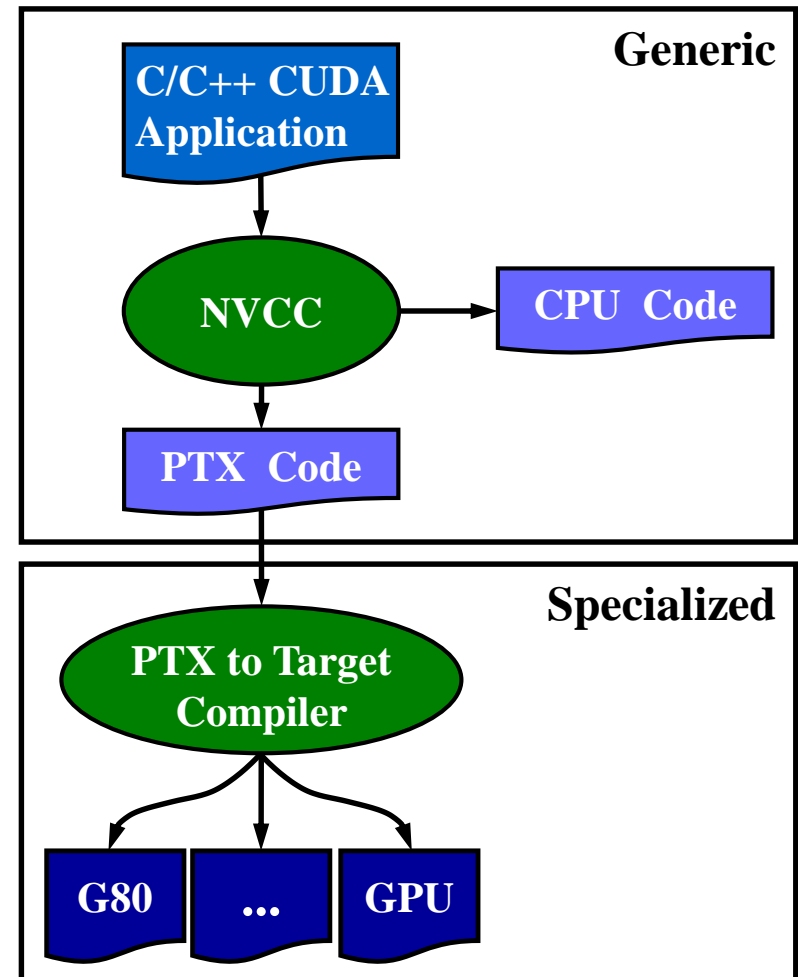
cudaMemcpyKind	Meaning	dst	src
cudaMemcpyHostToHost	Host → Host	host	host
cudaMemcpyHostToDevice	Host → Device	device	host
cudaMemcpyDeviceToHost	Device → Host	host	device
cudaMemcpyDeviceToDevice	Device → Device	device	device

host to host has the same effect as memcpy()



Program Compilation

- Any source file containing CUDA language must be compiled with NVCC
 - NVCC separates code running on the host from code running on the device
- Two-stage compilation:
 - Virtual ISA (Instruction Set Architecture)
 - PTX: Parallel Threads eXecutions
 - Device-specific binary object



Contents

- 1. Programming Model**
- 2. CUDA Language**
- 3. Example Code Study**
- 4. CPU & GPU Synchronization**
- 5. Multi-GPU**
- 6. Dynamic Parallelism**



Example 1: Hello World!

```
__global__ void mykernel(void) {  
  
}  
  
int main(void) {  
    mykernel<<<1,1>>>();  
    printf("Hello World!¥n");  
    return 0;  
}
```

■ Two new syntactic elements...

1. `__global__` indicates a function that runs on the device and is called from host code

2. `mykernel<<<1, 1>>>()`;

Triple angle brackets mark a call from host code to device code, which is called a “kernel launch”.



Example 2: Add 2 Numbers

```
__global__ void add(int *a, int *b, int *c) {  
    *c = *a + *b;  
}  
  
int main(void) {  
    int ha = 1, hb = 2, hc;  
    add<<<1,1>>>(&ha, &hb, &hc);  
    printf("c = %d\n", hc);  
    return 0;  
}
```



- This does not work!!
- `int ha, hb, hc` are in the host memory (DRAM), which cannot be used by device (GPU).
- We need to allocate variables in “device memory”.



The Correct main()

```
int main(void) {  
    int a=1, b=2, c;           // a, b, c的主机端备份  
    int *d_a, *d_b, *d_c;     // a, b, c的设备端备份  
    // 为变量a, b, c分配设备端内存  
    cudaMalloc((void **) &d_a, sizeof(int));  
    cudaMalloc((void **) &d_b, sizeof(int));  
    cudaMalloc((void **) &d_c, sizeof(int));  
    // 将输入拷贝到设备端  
    cudaMemcpy(d_a, &a, sizeof(int), cudaMemcpyHostToDevice);  
    cudaMemcpy(d_b, &b, sizeof(int), cudaMemcpyHostToDevice);  
    // 在GPU上启动核函数add()  
    add<<<1,1>>>(d_a, d_b, d_c);  
    // 将结果拷贝回主机端  
    cudaMemcpy(&c, d_c, sizeof(int), cudaMemcpyDeviceToHost);  
    // 释放内存  
    cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);  
    return 0;  
}
```

Example 3: Add 2 Vectors

- Let's first look at the sequential code!

```
// function definition
void VecAdd(int N, float* A, float* B, float* C)
{
    for (int i = 0; i < N; i++)
        C[i] = A[i] + B[i];
}

int main()
{
    ...
    VecAdd(N, Ah, Bh, Ch);
    ...
}
```



Parallel CUDA Code

- Use `blockIdx.x` as the index of the arrays
- Each thread processes 1 addition, for the elements indexed at `blockIdx.x`.

```
// Kernel definition
__global__ void VecAdd(float* A, float* B, float* C)
{
    int i = threadIdx.x;
    C[i] = A[i] + B[i];
}

int main()
{
    ...
    // Kernel invocation with N threads
    VecAdd<<<1, N>>>(Ad, Bd, Cd);
    ...
}
```

Alternative Implementation

- Using parallel thread instead

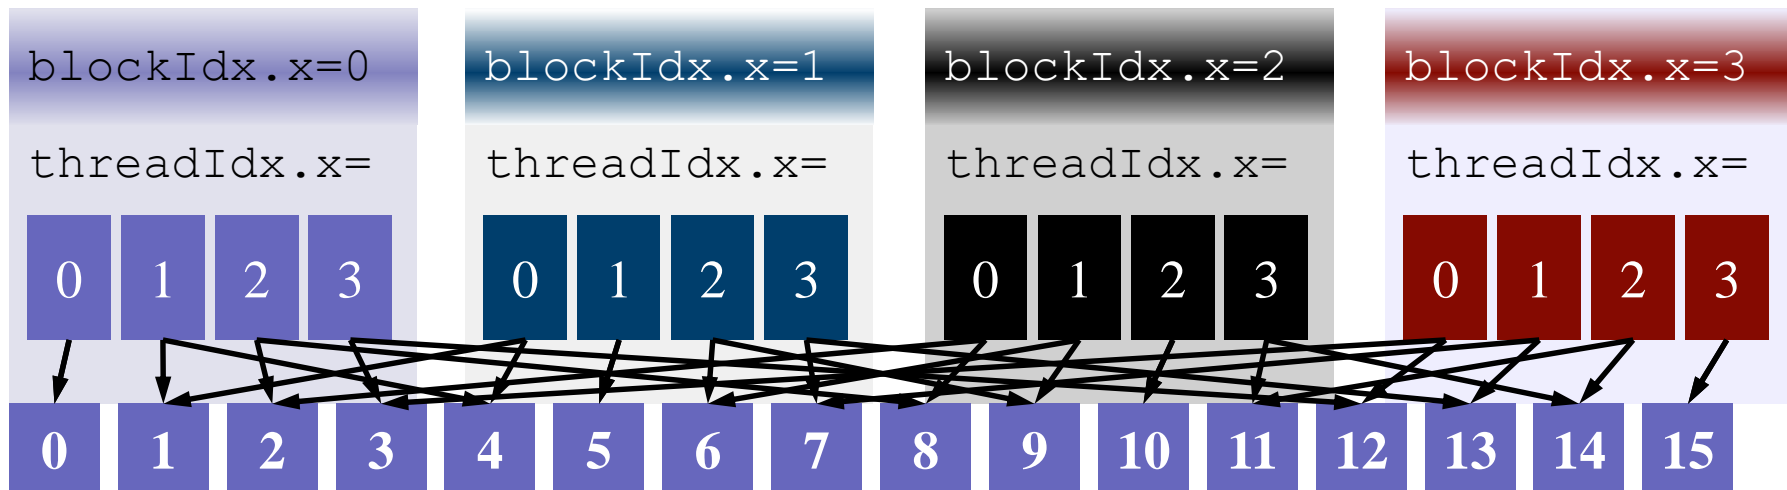
```
__global__ void add(int *a, int *b, int *c) {  
    c[blockIdx.x] = a[blockIdx.x] + b[blockIdx.x];  
}  
  
int main(void) {  
    int a[N], b[N], c[N];  
    int *d_a, *d_b, *d_c;  
    add<<<N, 1>>>>(d_a, d_b, d_c);  
    ...  
}
```

- N blocks and each block has 1 thread.
- Which one is better?
 - Threads in the same block can communicate, synchronize with others, but the number of threads per block is limited.



3rd Implementation

- Using multiple threads and multiple blocks
- Suppose $N=16$, grid size = 4, and block size = 4



- How to index 16 elements of an array?
 - Method 1: $\text{index} = \text{blockIdx.x} * 4 + \text{threadIdx.x}$
 - Method 2: $\text{index} = \text{threadIdx.x} * 4 + \text{blockIdx.x}$



The General Case

- Use the built-in variable `blockDim.x` for threads per block.

```
__global__ void add(int *a, int *b, int *c) {  
    int index = threadIdx.x + blockIdx.x * blockDim.x;  
    c[index] = a[index] + b[index];  
}  
  
int main(void) {  
    int a[N], b[N], c[N];  
    int *d_a, *d_b, *d_c;  
    ...  
    add<<<N/BS, BS>>>>(d_a, d_b, d_c);  
    ...  
}
```

What if N is not a multiple of BS?

- BS is block size (number of threads per block)



An Even More General Case

```
__global__ void add(int *a, int *b, int *c, int n) {  
    int index = threadIdx.x + blockIdx.x * blockDim.x;  
    if (index < n)  
        c[index] = a[index] + b[index];  
}  
  
int main(void) {  
    int a[N], b[N], c[N];  
    int *d_a, *d_b, *d_c;  
    ...  
    add<<< (N + BS - 1) / BS, BS >>> (d_a, d_b, d_c, N);  
    ...  
}
```

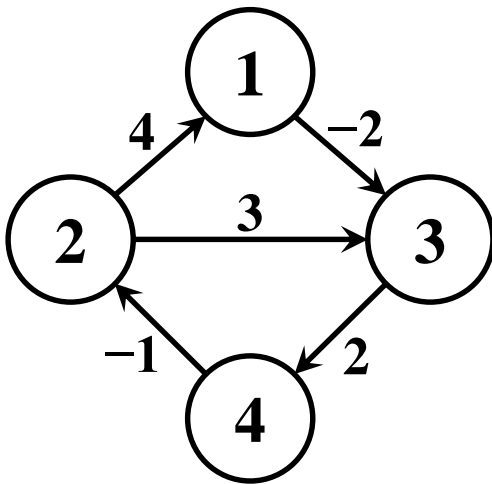
- The kernel function can have branches, but with a price to pay...



Example4: APSP (All Pair Shortest Paths)

- Given a weighted directed graph $G(V, E, W)$, where $|V| = n$, $|W|=m$, and $W>0$, find the shortest path of all pairs of vertices (v_i, v_j) .

- Example:



0	INF	-2	INF
4	0	3	INF
INF	INF	0	2
INF	-1	INF	0

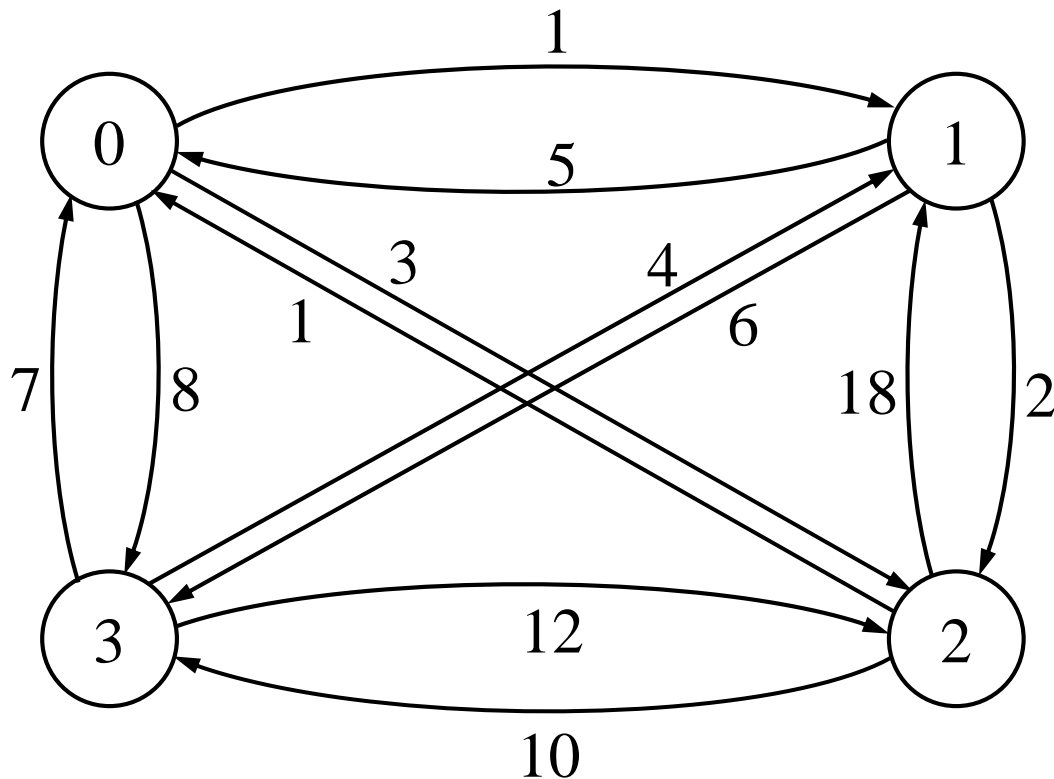
Initial
weight

0	-1	-2	0
4	0	2	4
5	1	0	2
3	-1	1	0

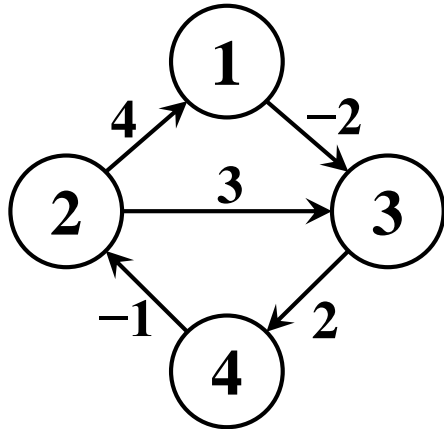
Final
result



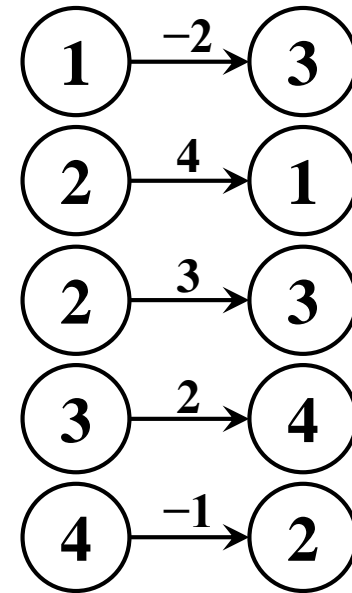
A Four-City TSP

$$\begin{bmatrix} 0 & 1 & 3 & 8 \\ 5 & 0 & 2 & 6 \\ 1 & 18 & 0 & 10 \\ 7 & 4 & 12 & 0 \end{bmatrix}$$


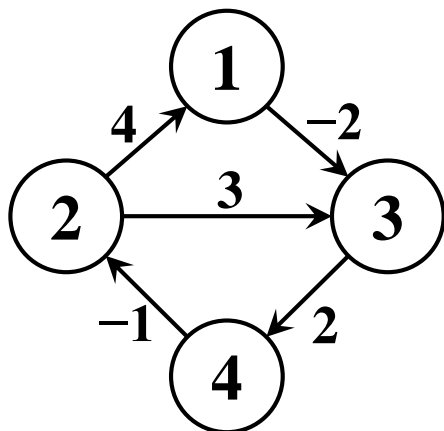
有向非完全图——求最短路径APSP (All Pair Shortest Paths)



	1	2	3	4
1	0	∞	-2	∞
2	4	0	3	∞
3	∞	∞	0	2
4	∞	-1	∞	0



假设所有路径都必须经过第1点时 ($k = 1$)



//经过1号顶点

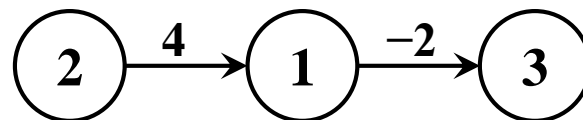
```
for(i=1;i<=n;i++)
```

```
    for(j=1;j<=n;j++)
```

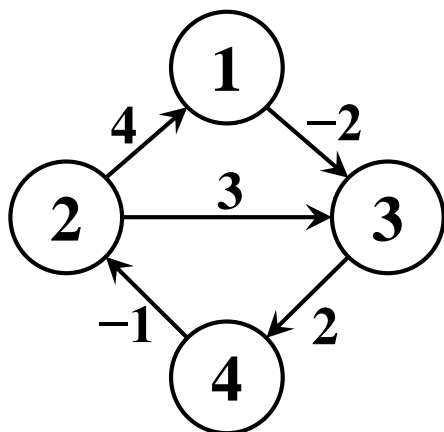
```
        if(e[i][j]>e[i][1]+e[1][j])
```

```
            e[i][j]=e[i][1]+e[1][j];
```

	1	2	3	4
1	0	∞	-2	∞
2	4	0	2	∞
3	∞	∞	0	2
4	∞	-1	∞	0



假设所有路径都必须经过第2点时 ($k = 2$)



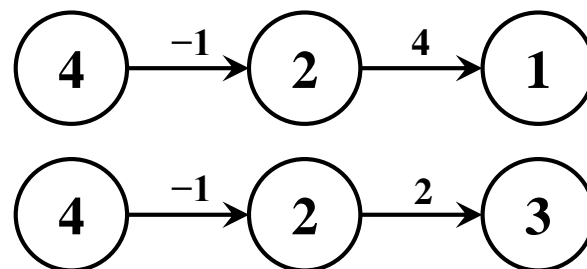
//经过1号顶点

```
for(i=1;i<=n;i++)
    for(j=1;j<=n;j++)
        if(e[i][j]>e[i][1]+e[1][j])
            e[i][j]=e[i][1]+e[1][j];
```

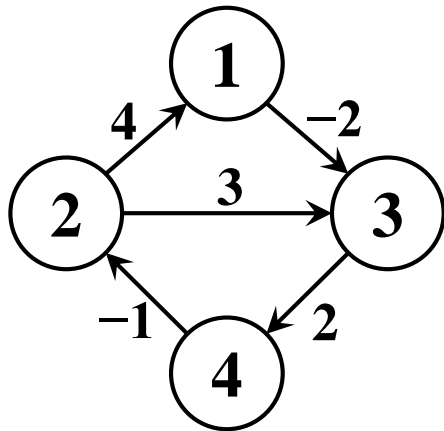
//经过2号顶点

```
for(i=1;i<=n;i++)
    for(j=1;j<=n;j++)
        if(e[i][j]>e[i][2]+e[2][j])
            e[i][j]=e[i][2]+e[2][j];
```

	1	2	3	4
1	0	∞	-2	∞
2	4	0	2	∞
3	∞	∞	0	2
4	3	-1	1	0



假设所有路径都必须经过第3点时 ($k = 3$)



	1	2	3	4
1	0	∞	-2	0
2	4	0	2	4
3	∞	∞	0	2
4	3	-1	1	0

//经过1号顶点

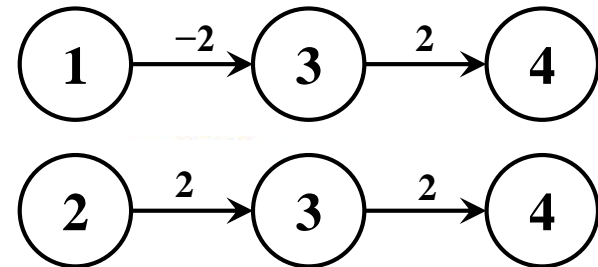
```
for(i=1;i<=n;i++)
    for(j=1;j<=n;j++)
        if(e[i][j]>e[i][1]+e[1][j])
            e[i][j]=e[i][1]+e[1][j];
```

//经过2号顶点

```
for(i=1;i<=n;i++)
    for(j=1;j<=n;j++)
        if(e[i][j]>e[i][2]+e[2][j])
            e[i][j]=e[i][2]+e[2][j];
```

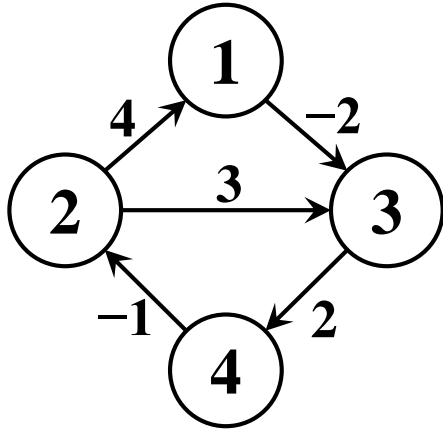
//经过3号顶点

```
for(i=1;i<=n;i++)
    for(j=1;j<=n;j++)
        if(e[i][j]>e[i][3]+e[3][j])
            e[i][j]=e[i][3]+e[3][j];
```



假设所有路径都必须经过第4点时 ($k = 4$)

How to parallelize it?



// 弗洛伊德算法串行代码

```
for(k=1;k<=n;k++)
```

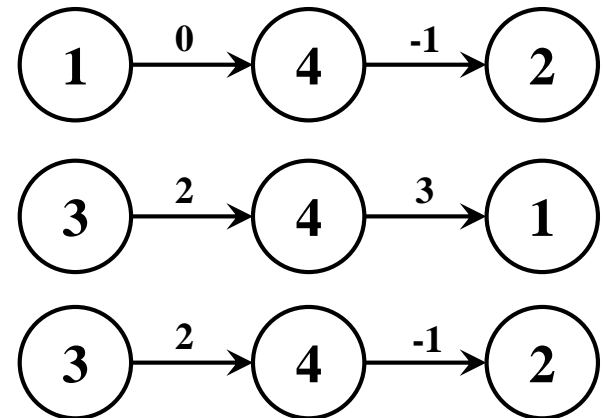
```
    for(i=1;i<=n;i++)
```

```
        for(j=1;j<=n;j++)
```

```
            if(e[i][j]>e[i][k]+e[k][j])
```

```
                e[i][j]=e[i][k]+e[k][j];
```

	1	2	3	4
1	0	-1	-2	0
2	4	0	2	4
3	5	1	0	2
4	3	-1	1	0



Implementation 1

- 1 block and n threads
- Thread i updates the SP for vertex i.

```
__global__ void FW_APSP(int k, int D[n][n]) {  
    int i = threadIdx.x;  
    for (int j = 0; j < n; j++)  
        if (D[i][j] > D[i][k] + D[k][j])  
            D[i][j] = D[i][k] + D[k][j];  
}  
  
int main() {  
    ...  
    for (int k = 0; k < n, k++)  
        FW_APSP<<<1, n>>>>(k, D);  
}
```

Simple! But can it be faster ?



Implementation 2

- Each thread updates one pair of vertices
 - Increase parallelism from n to n^2

```
__global__ void FW_APSP(int k, int D[n][n]) {  
    int i = threadIdx.x;  
    int j = threadIdx.y;  
    if (D[i][j] > D[i][k] + D[k][j])  
        D[i][j] = D[i][k] + D[k][j];  
}  
  
int main() {  
    ...  
    dim3 threadsPerBlock(n, n);  
    for (int k = 0; k < n; k++)  
        FW_APSP<<<1, threadsPerBlock>>>(k, D);  
}
```

How about the for-loop of k ?



Implementation 3

```
__global__ void FW_APSP(int D[n][n]) {
    int i = threadIdx.x;
    int j = threadIdx.y;
    for (int k = 0; k < n, k++)
        if (D[i][j] > D[i][k] + D[k][j])
            D[i][j] = D[i][k] + D[k][j];
}

int main() {
    ...
    dim3 threadsPerBlock(n, n);
    FW_APSP<<<1, threadsPerBlock>>>(D);
}
```



- It is a synchronous computation
 - There is data dependency on k...



Add __syncthreads ()

```
__global__ void FW_APSP(int D[n][n]) {  
    int i = threadIdx.x;  
    int j = threadIdx.y;  
    for (int k = 0; k < n, k++) {  
        if (D[i][j] > D[i][k] + D[k][j])  
            D[i][j] = D[i][k] + D[k][j];  
        __syncthreads ();  
    }  
}  
  
int main() {  
    ...  
    dim3 threadsPerBlock(n, n);  
    FW_APSP<<<1, threadsPerBlock>>>(D);  
}
```



Contents

- 1. Programming Model**
- 2. CUDA Language**
- 3. Example Code Study**
- 4. CPU & GPU Synchronization**
- 5. Multi-GPU**
- 6. Dynamic Parallelism**



Asynchronous Functions

- To facilitate concurrent execution between host and device, most CUDA function calls are **asynchronous**:
 - Control is returned to the host thread before the device has completed the requested task.
 - **But function calls from a kernel are serialized on GPU**
- Asynchronous functions:
 - Kernel launches
 - Asynchronous memory copy and set options:
cudaMemcpyAsync, cudaMemcpyAsync
 - cudaMemcpy within the same device
 - H2D cudaMemcpy of 64kB or less
 - cudaEvent functions

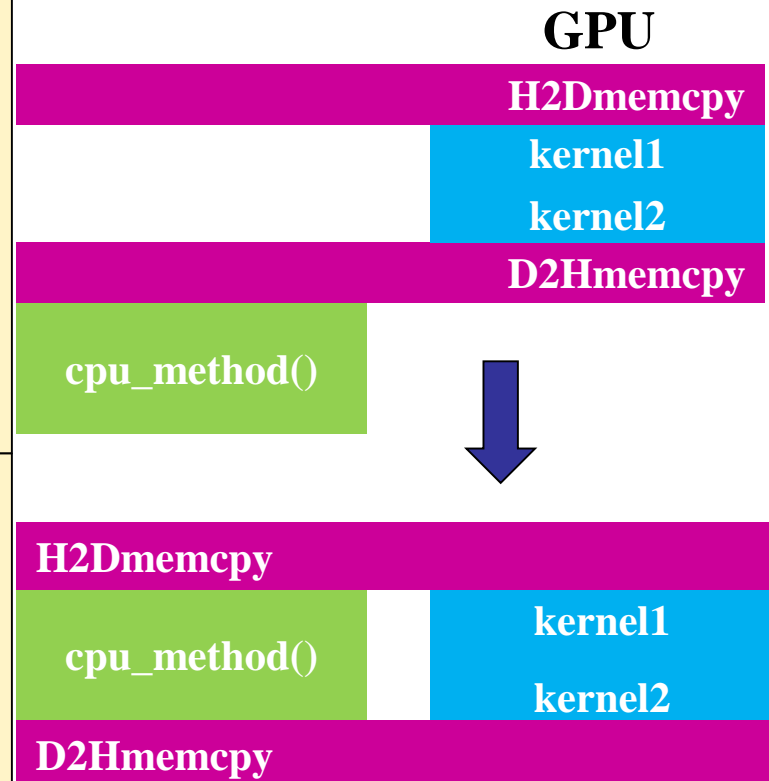


Why Use Asynchronous Functions?

- Overlap CPU computation with the GPU computation or data transfer

```
void main() {  
    cudaMemcpy(/**/, H2D);  
    kernel1<<<grid, block>>>();  
    kernel2<<<grid, block>>>();  
    cudaMemcpy(/**/, D2H);  
    cpu_method();  
}
```

```
void main() {  
    cudaMemcpy(/**/, H2D);  
    kernel1<<<grid, block>>>();  
    kernel2<<<grid, block>>>();  
    cudaMemcpyAsync(/**/, D2H);  
    cpu_method();  
}
```



Risk of Using Asynchronous Functions

- Programmer must enforce synchronization between GPU and CPU when there is **data dependency**

```
void main() {  
    cudaMemcpyAsync(d_a, h_a, count, H2D);  
    kernel<<<grid, block>>>(d_a);  
    cudaMemcpyAsync(h_a, d_a, count, D2H);  
    cpu_method(h_a);  
}
```

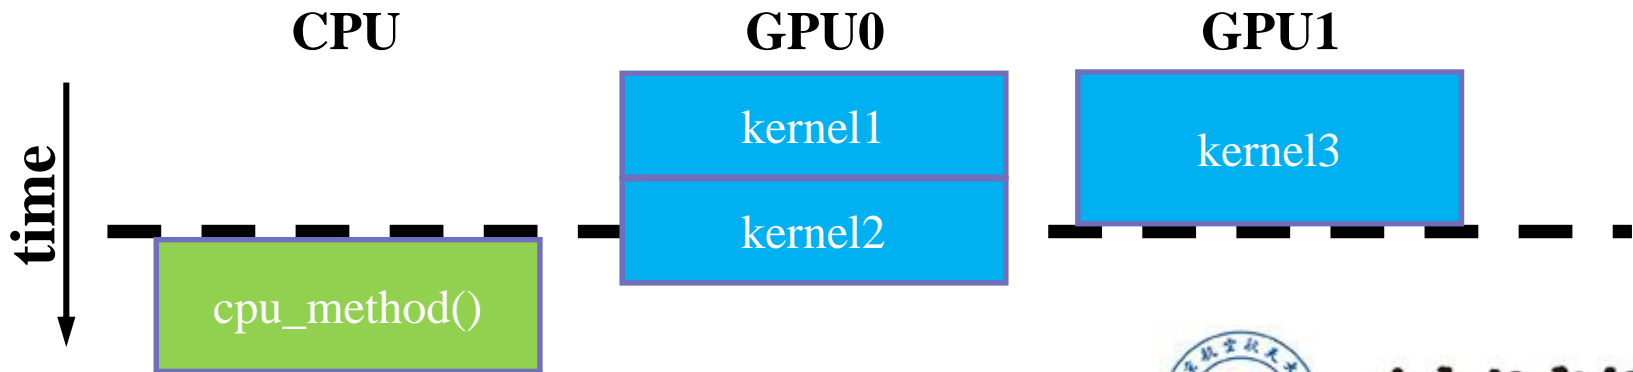


Synchronization between CPU & GPU

- Device based: `cudaDeviceSynchronize()`
 - Block a **CPU** thread until all issued CUDA calls to a device complete
- Context based: `cudaThreadSynchronize()`
 - Block a **CPU** thread until all issued CUDA calls from the thread complete
- Stream based: `cudaStreamSynchronize(stream-id)`
 - Block a **CPU** thread until all CUDA calls in stream stream-id complete
- Event based:
 - `cudaEventSynchronize(event)`
 - Block a **CPU** thread until event is recorded
 - `cudaStreamWaitEvent(stream-id, event)`
 - Block a **GPU** stream until event reports completion

Device Synchronization Example

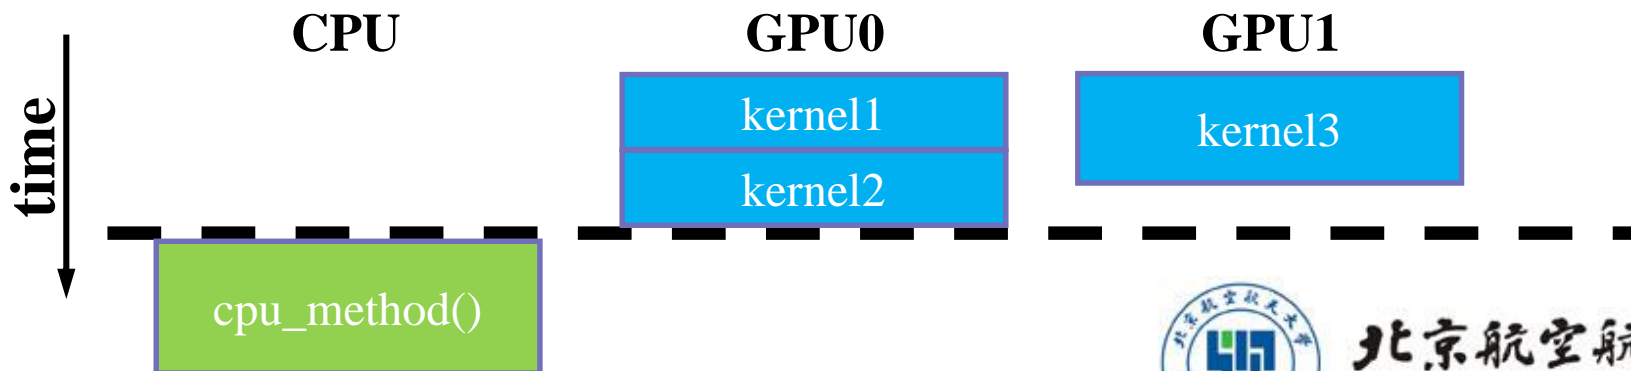
```
void main() {  
    cudaSetDevice(0);  
    kernel1<<<grid, block>>>();  
    kernel2<<<grid, block>>>();  
    cudaSetDevice(1);  
    kernel3<<<grid, block>>>();  
    cudaDeviceSynchronize();  
    cpu_method();  
}
```



Thread Synchronization Example

```
void main() {  
    cudaSetDevice(0);  
    kernel1<<<grid, block>>>();  
    kernel2<<<grid, block>>>();  
    cudaSetDevice(1);  
    kernel3<<<grid, block>>>();  
    cudaThreadSynchronize();  
    cpu_method();  
}
```

deprecated



CUDA Event

CUDA中Event主要用于在流的执行中添加标记点，用于检查正在执行的流是否到达给定点。

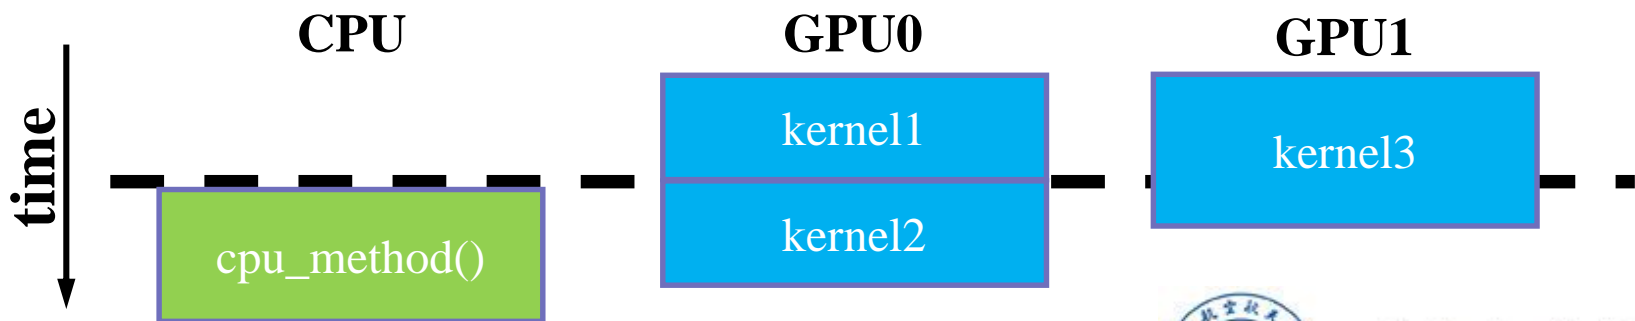
Data type: `cudaEvent_t`

- `cudaError_t cudaEventCreate(cudaEvent_t* event)`
 - Create CUDA event
- `cudaError_t cudaEventRecord(cudaEvent_t event, cudaStream_t stream = 0)`
 - Record CUDA event
 - If stream is non-zero, the event is recorded after all preceding operations in the stream have been completed
 - Since operation is asynchronous, `cudaEventQuery()` and/or `cudaEventSynchronize()` must be used to determine when the event has actually been recorded
- `cudaError_t cudaEventSynchronize(cudaEvent_t event)`
 - Wait until the completion of all device work preceding the most recent call to `cudaEventRecord()`



Event Synchronization Example

```
void main() {  
    cudaSetDevice(0);  
    kernel1<<<grid, block>>>();  
    cudaEventRecord(event)  
    kernel2<<<grid, block>>>();  
    cudaSetDevice(1);  
    kernel3<<<grid, block>>>();  
    cudaEventSynchronize(event)  
    cpu_method();  
}
```



Kernel Time Measurement Example

```
1  cudaEvent_t start, stop;
2  cudaEventCreate(&start);
3  cudaEventCreate(&stop);
4  cudaEventRecord(start);
5  kernel<<<block, thread>>>();
6  cudaEventRecord(stop);
7  cudaEventSynchronize(stop);
8  float time;
9  cudaEventElapsedTime(&time, start, stop);
```

Create event

Record event

Record event and
synchronize

Compute the event duration



Contents

- 1. Programming Model**
- 2. CUDA Language**
- 3. Example Code Study**
- 4. CPU & GPU Synchronization**
- 5. Multi-GPU**
- 6. Dynamic Parallelism**



Multi-GPUs

- Multi-GPUs within a node
 - A single CPU thread, multiple GPU
 - Multiple CPU threads belonging to the same process, such as pthread or openMP
- Multiple GPUs on multiple nodes
 - Need to go through network API, such as MPI



Single Thread Multi-GPUs

- All CUDA calls are issued to the current GPU
 - **cudaSetDevice()** sets the current GPU

```
// Run independent kernel on each CUDA device
int numDevs = 0;
cudaGetNumDevices (&numDevs);
for (int d = 0; d < numDevs; d++) {
    cudaSetDevice (d);
    kernel<<<blocks, threads>>>(args);
}
```

- Asynchronous calls (kernels, memcopies) don't block switching the GPU

```
cudaSetDevice(0);
kernel<<<...>>> (...);
cudaSetDevice(1);
kernel<<<...>>> (...);
```



Using CUDA with OpenMP

- Put CUDA functions inside the parallel region
- General setting:
 - The number of CPU threads is the same as the number of CUDA devices.
 - Each CPU thread controls a different device, processing its portion of the data.
- It's possible to use more CPU threads than there are CUDA devices.
 - Several CPU threads will be allocating resources and launching kernel on the same device, which will slow down the performance.



Example: cudaOMP.cu

```
...
cudaGetDeviceCount (&num_gpus);
...
omp_set_num_threads (num_gpus);

// 生成和CUDA devices同等数量的CPU线程
#pragma omp parallel
{
    unsigned int cpu_thread_id = omp_get_thread_num();
    unsigned int num_cpu_threads = omp_get_num_threads();
    cudaSetDevice (cpu_thread_id);
    int gpu_id = -1;
    cudaGetDevice (&gpu_id);
    printf("CPU thread %d (of %d) uses CUDA device %d\n",
          cpu_thread_id, num_cpu_threads, gpu_id);
    ...
}
```



Using CUDA with MPI

```
int main(int argc, char* argv[]) {
    int my_rank, size;
    int A[32];
    int i;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    printf("I am %d of %d\n", my_rank, size);
    for(i = 0; i < 32; i++)
        A[i] = my_rank + 1;
    launch(A); // a call to launch CUDA kernel
    MPI_Barrier(MPI_COMM_WORLD);
    MPI_Finalize();
    return 0;
}
```

Example: launch(A)

```
extern "C"  
void launch(int *A) {  
    int *dA;  
    cudaMalloc((void**) &dA, sizeof(int) * 32);  
    cudaMemcpy(dA, A, sizeof(int) * 32,  
               cudaMemcpyHostToDevice);  
    kernel<<<1, 32>>>(dA);  
    cudaMemcpy(A, dA, sizeof(int) * 32,  
               cudaMemcpyDeviceToHost);  
    cudaFree(dA);  
}
```



Sharing Data between GPUs

Options

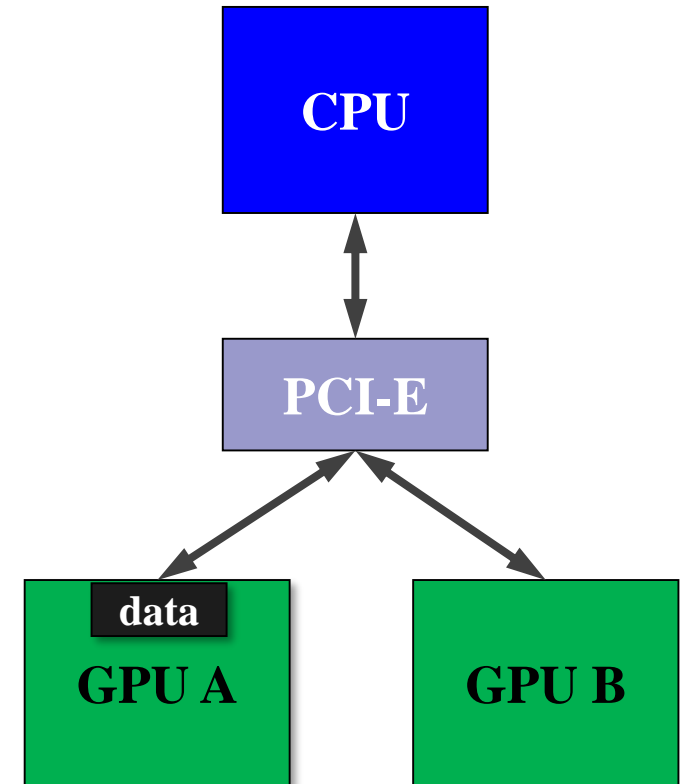
1. Explicit copies via host
2. Zero-copy shared host array
3. Peer-to-peer memory copy



1. Explicit Copies via Host

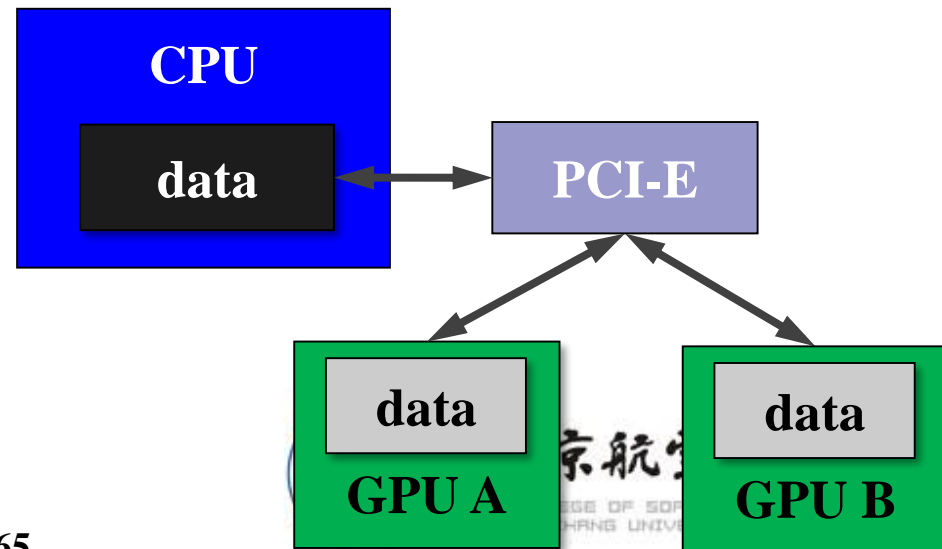
- CPU explicitly copies data from device A to device B
- Example

```
cudaSetDevice(0);  
cudaMemcpy(DM1, HM, n, D2H);  
cudaSetDevice(1);  
cudaMemcpy(HM, DM2, n, H2D);
```



2. Using Zero-copy

- Zero-copy” refers to direct device access to host memory
- Device threads can read directly from host memory over PCI-E without using cudaMemcpy H2D or D2H
- The host memory must be pinned (page-locked)
 - Pageable memory cannot be directly accessed by the GPU because of the OS virtual memory mechanism.



Host Memory Allocation

- `malloc()`
 - Regular C library memory allocation
 - Managed by host
 - Cannot be directly accessed by GPUs
 - Must be copied to device memory via `cudaMemcpy()`
 - `cudaMallocHost(void** hostPtr, size_t size)`
 - Allocate pinned (page-locked) host memory for higher `cudaMemcpy` performance
 - Used with `cudaMemAsync()` for async memory copy or CUDA stream
 - `cudaHostAlloc(void** hostPtr, size_t size, unsigned int flags)`
 - Add the flag “`cudaHostAllocMapped`” to allocate pinned host memory for higher `cudaMemcpy` performance
 - Add the flag “`cudaHostAllocPortable`” to allocate shared host memory for “Zero copy”
- ➔ **All these functions return host memory pointer**

Zero Copy via `cudaHostAlloc()`

- Allocate **host memory** using `cudaHostAlloc()`
 - It returns a host pointer.
 - It can enable faster host memory access.
 - Must add flag `cudaHostAllocMapped` for **page-locked**
 - Add flag `cudaHostAllocPortable` for **sharing among all devices**
- Bind host pointer to device pointer using
`cudaHostGetDevicePointer(void**, void*, 0)`

```
int *hostPtr, *dev0Ptr, *dev1Ptr;
cudaHostAlloc(&hostPtr, 10, cudaHostAllocMapped|
cudaHostAllocPortable);
cudaSetDevice(0);
cudaHostGetDevicePointer(&dev0Ptr, hostPtr, 0);
kernel<<1,10>>(dev0Ptr);
cudaSetDevice(1);
cudaHostGetDevicePointer(&dev1Ptr, hostPtr, 0);
kernel<<1,10>>(dev1Ptr);
```



Pitfalls of using Zero-copy

- PCI-E is lower in bandwidth and higher in latency than GPU global memory
 - ➔ **Access speed is slower than global memory**
- Use zero-copy if:
 - The data is only accessed once or few times
 - Generate data on the device and copy back to host without reuse
 - Kernel(s) that access the memory are compute bound



3. Peer-to-Peer Memcpy

- Direct copy from pointer on GPU A to pointer on GPU B
- Using two functions

```
cudaError_t cudaMemcpyPeer(void *dst, int dstDevice,  
const void* src, int srcDevice, size_t count)
```

```
cudaError_t cudaMemcpyPeerAsync(void *dst, int dstDevice,  
const void* src, int srcDevice, size_t count,  
cuda_stream_t stream=0)
```



Example: P2P memcpy

```
cudaSetDevice(0);           // 将0号设备置为当前设备
float *p0;
size_t size = 1024 * sizeof(float);
cudaMalloc(&p0, size);      // 在0号设备上分配内存
cudaSetDevice(1);          // 将1号设备置为当前设备
float *p1;
cudaMalloc(&p1, size);      // 在1号设备上分配内存
cudaSetDevice(0);          // 将0号设备置回当前设备
Kernel1<<<1000, 128>>>(p0); // 在0号设备上启动核函数
cudaSetDevice(1);          // 将1号设备置回当前设备
cudaMemcpyPeer(p1, 1, p0, 0, size); // 将p0拷贝到p1
Kernel1<<<1000, 128>>>(p1); // 在1号设备上启动核函数
```



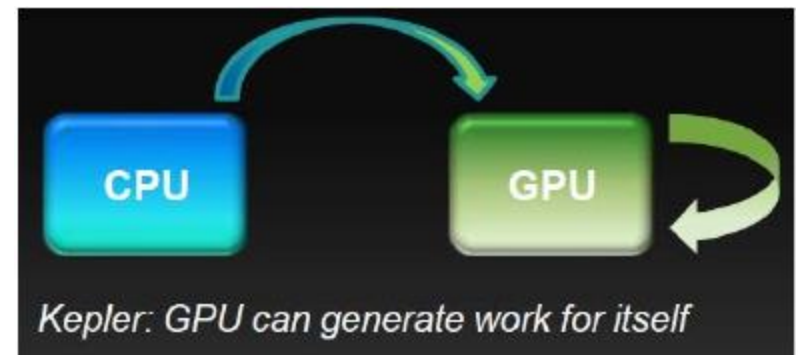
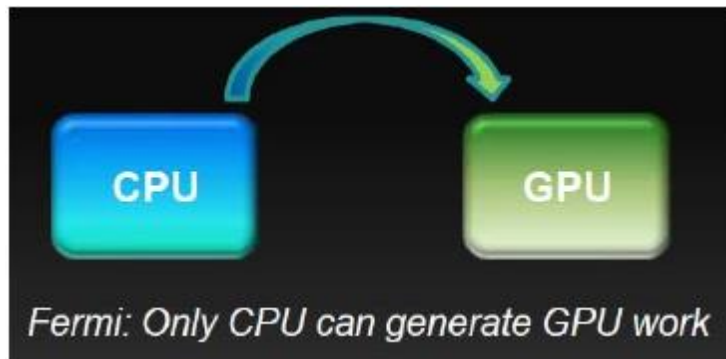
Contents

- 1. Programming Model**
- 2. CUDA Language**
- 3. Example Code Study**
- 4. CPU & GPU Synchronization**
- 5. Multi-GPU**
- 6. Dynamic Parallelism**



Dynamic Parallelism

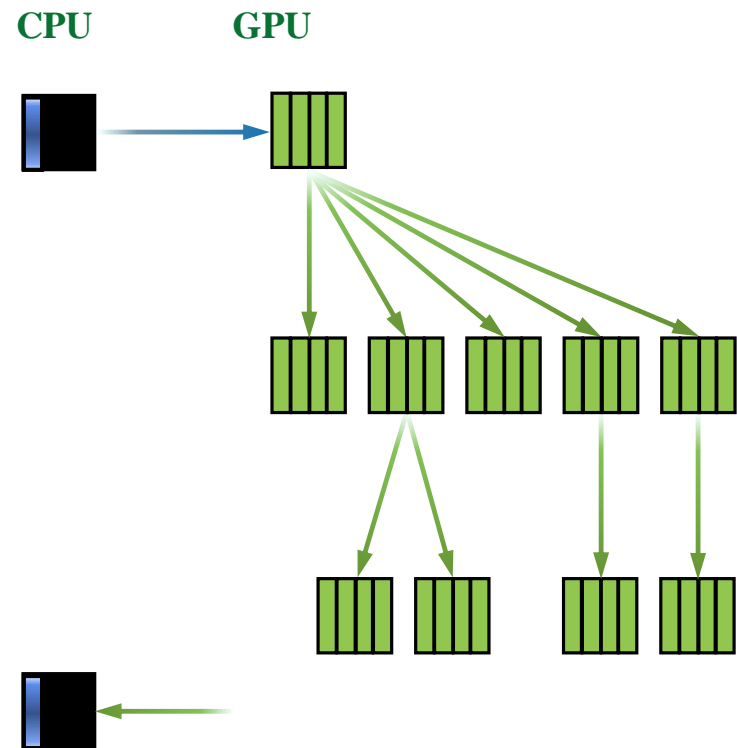
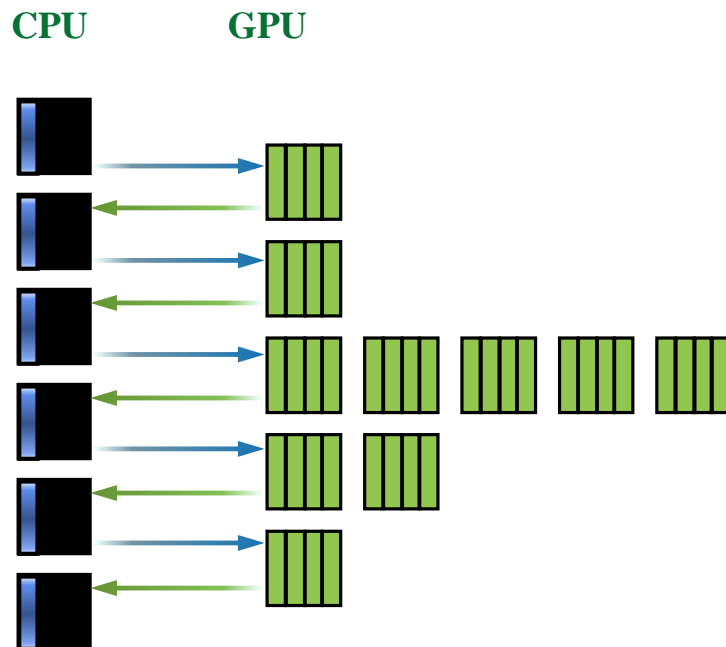
- The ability to launch new **grids** from the **GPU**
 - Dynamically
 - Simultaneously
 - Independently
- Supported from CUDA5.0 on devices of Compute Capability 3.5 or higher



What does It Mean?

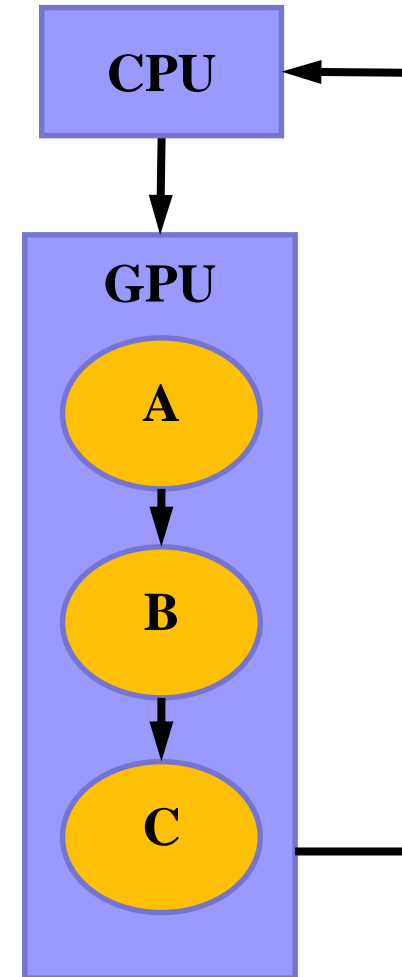
- Reduce the number of kernel launches

Dynamic Parallelism



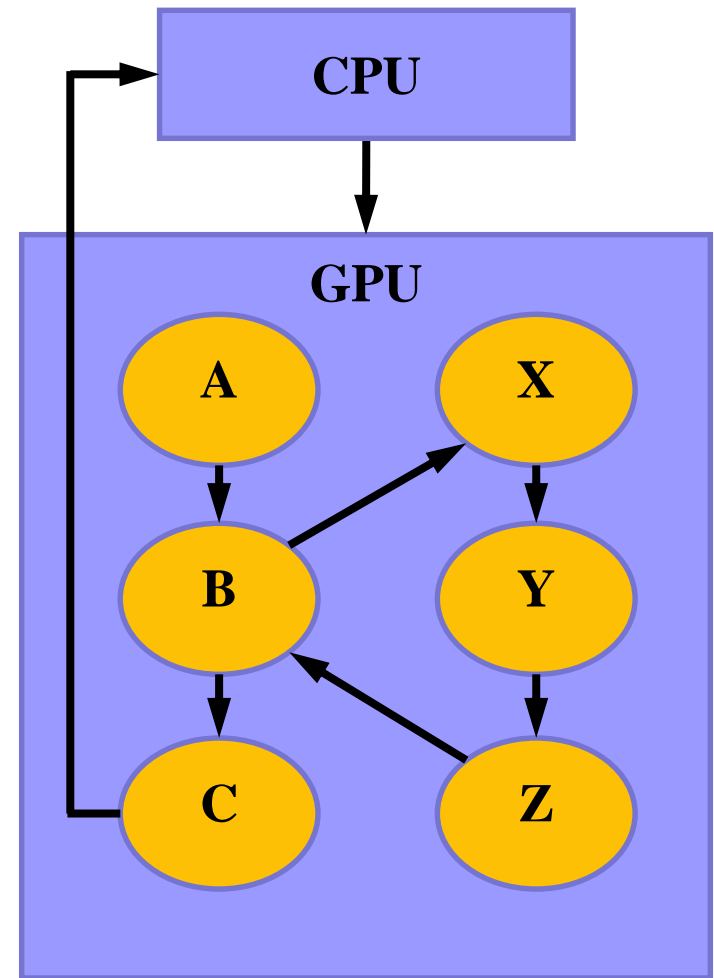
Dependency in CUDA

```
void main() {  
    float *data;  
    do_stuff(data);  
    A<<<...>>>(data);  
    B<<<...>>>(data);  
    C<<<...>>>(data);  
    cudaDeviceSynchronize();  
    do_more_stuff(data);  
}
```



Nested Dependency

```
void main() {  
    float *data;  
    do_stuff(data);  
    A<<<...>>>(data);  
    B<<<...>>>(data);  
    C<<<...>>>(data);  
    cudaDeviceSynchronize();  
    do_more_stuff(data);  
}  
  
__global__ void B(float *data) {  
    do_stuff(data);  
    X<<<...>>>(data);  
    Y<<<...>>>(data);  
    Z<<<...>>>(data);  
    cudaDeviceSynchronize();  
    do_more_stuff(data);  
}
```



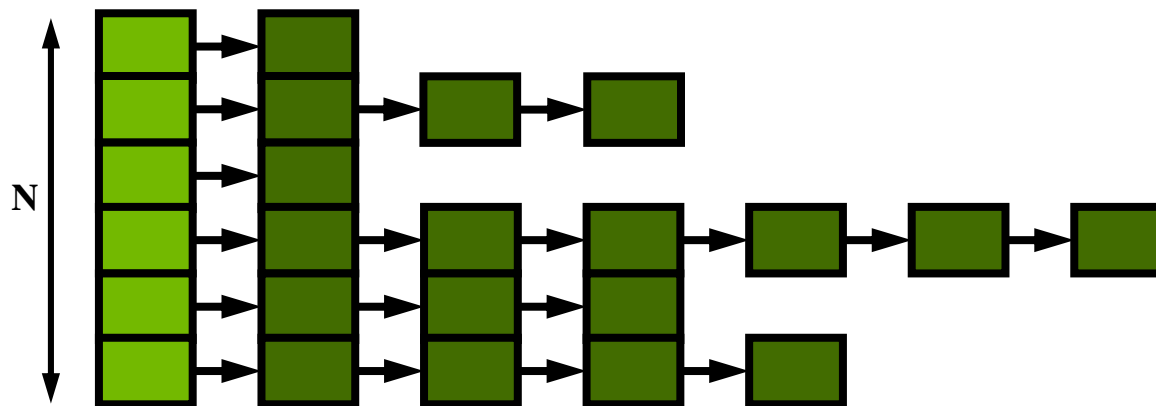
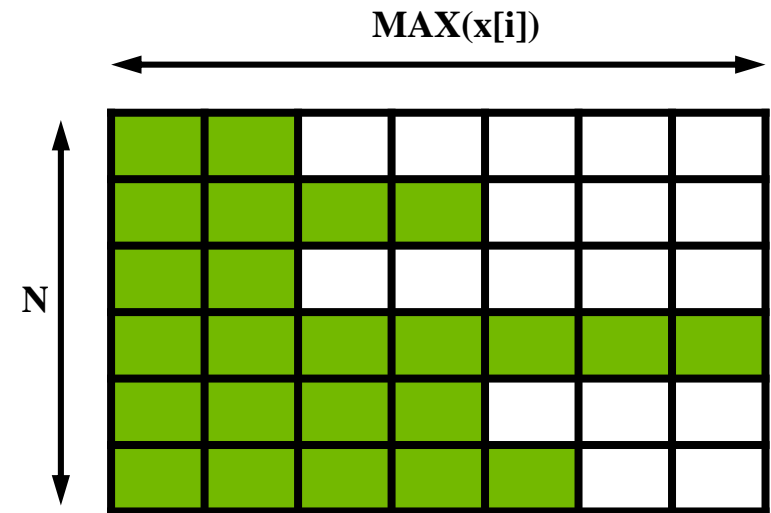
What is DP Good for?

- Dynamic block size and grid size
- Dynamic work generation
- Nested parallelism
- Library calls
- Parallel recursion



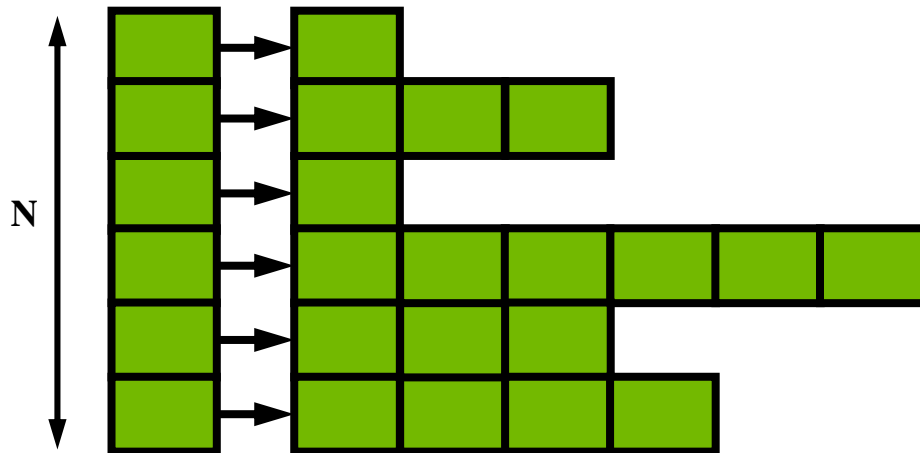
1. Dynamic Block Size and Grid Size

```
for i = 1 to N
  for j = 1 to x[i]
    convolution(i, j)
  next j
next i
```

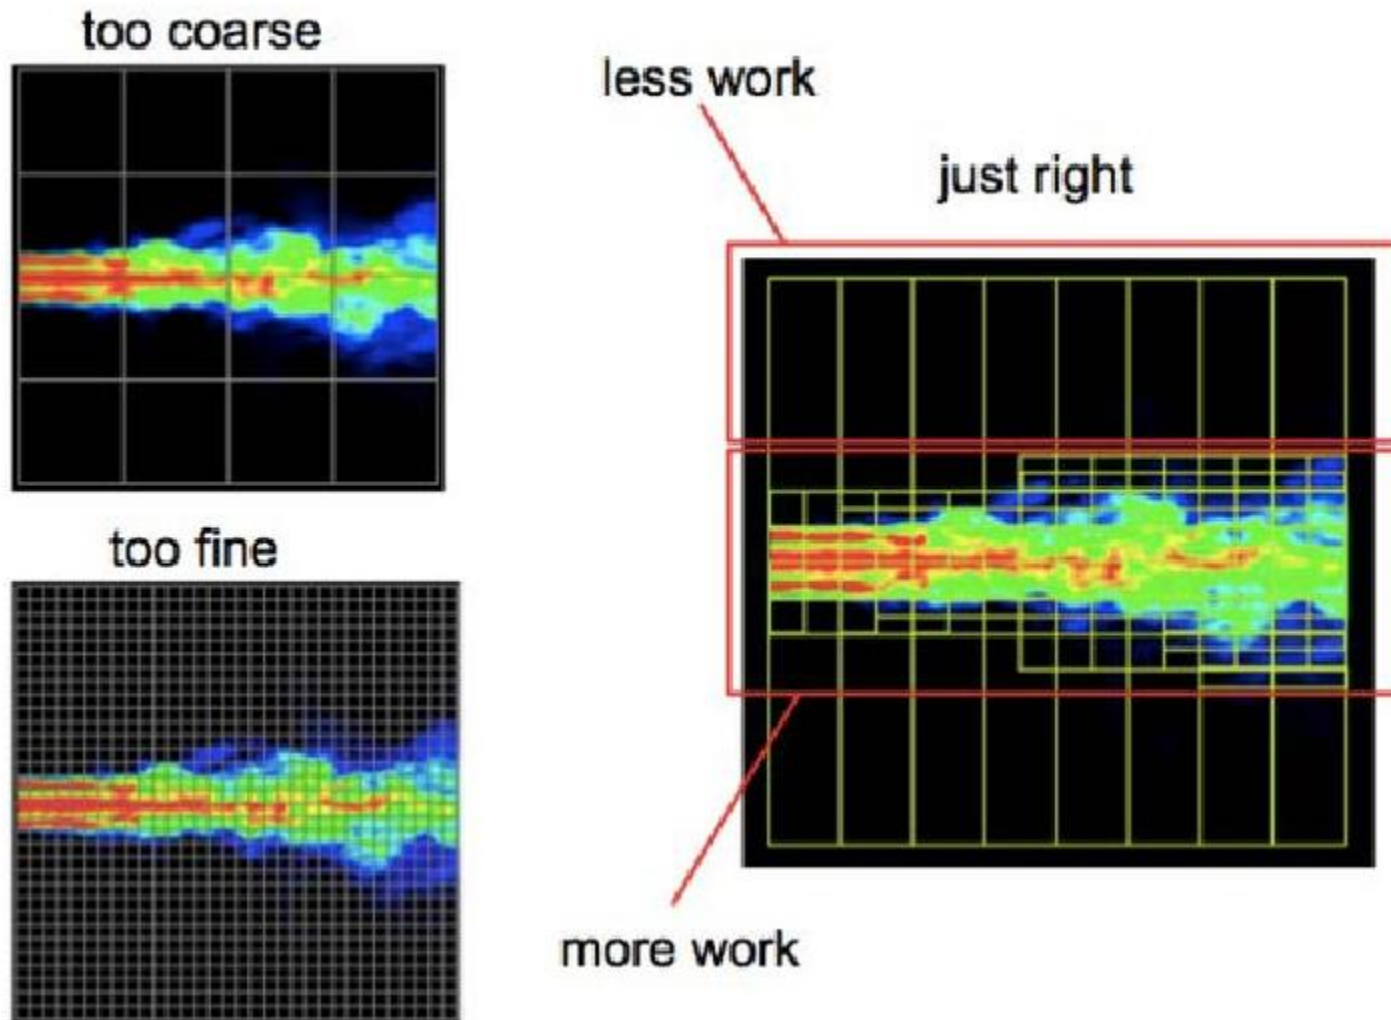


1. Dynamic Block Size with DP

```
__global__ void convolution(int x[]) {  
    for j = 1 to x[blockIdx]  
        kernel<<<...>>>(blockIdx, j)  
    }  
    ...  
    convolution<<<N, 1>>>(x);  
}
```

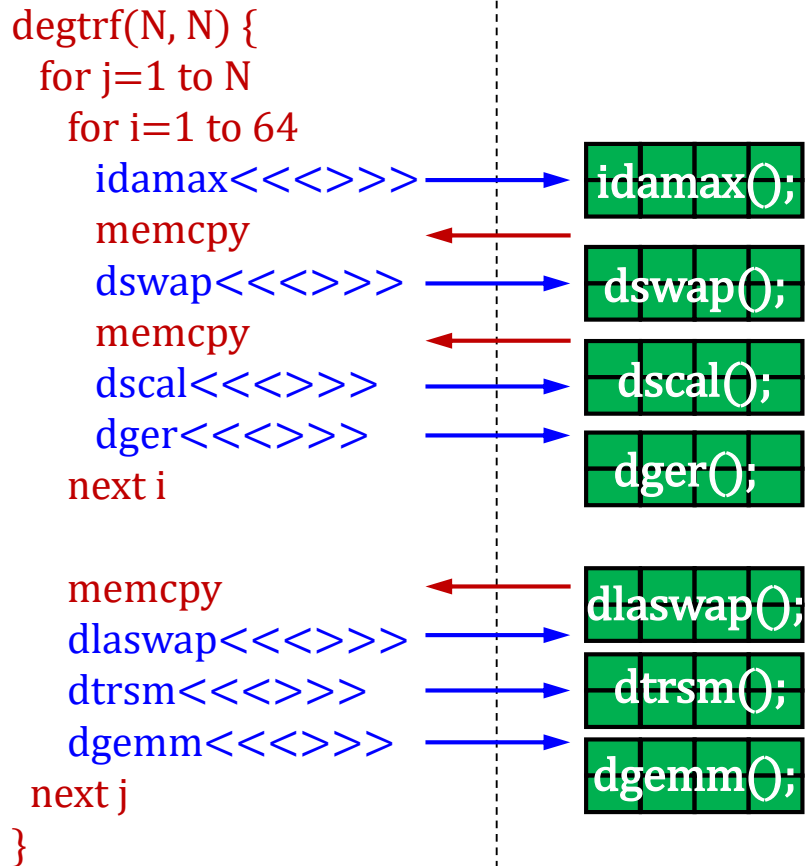


2. Dynamic Work Generation



3. Nested Parallelism

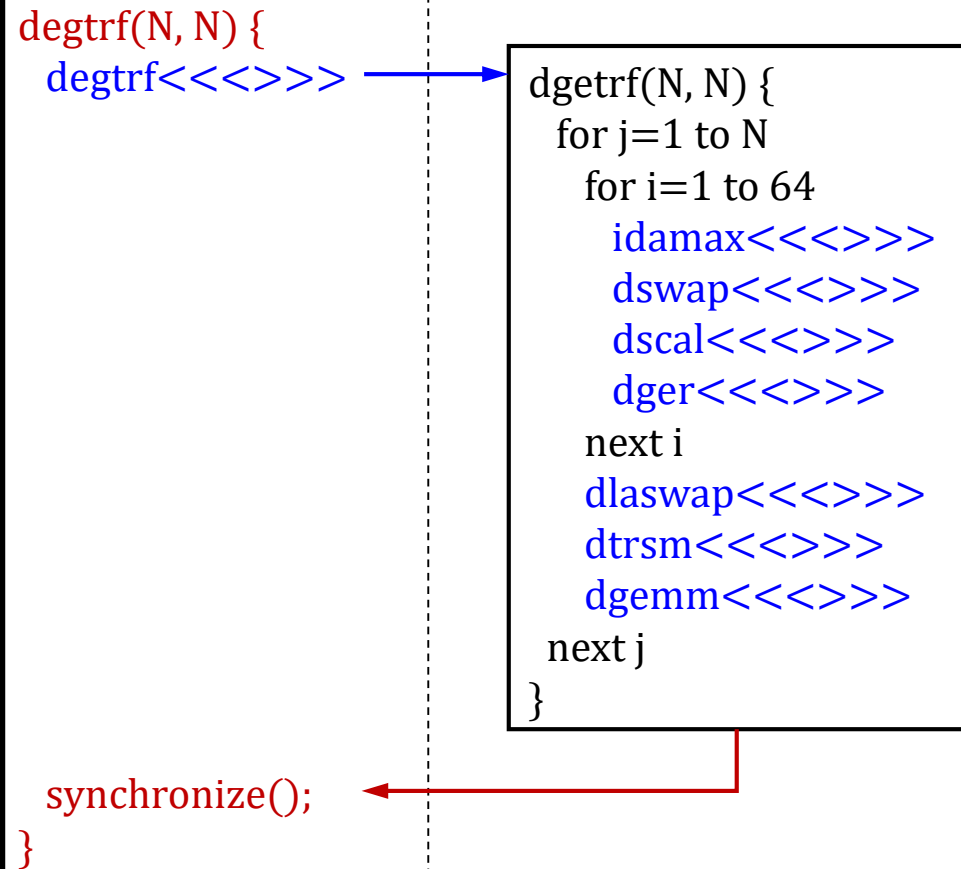
LU decomposition (Fermi)



CPU Code

GPU Code

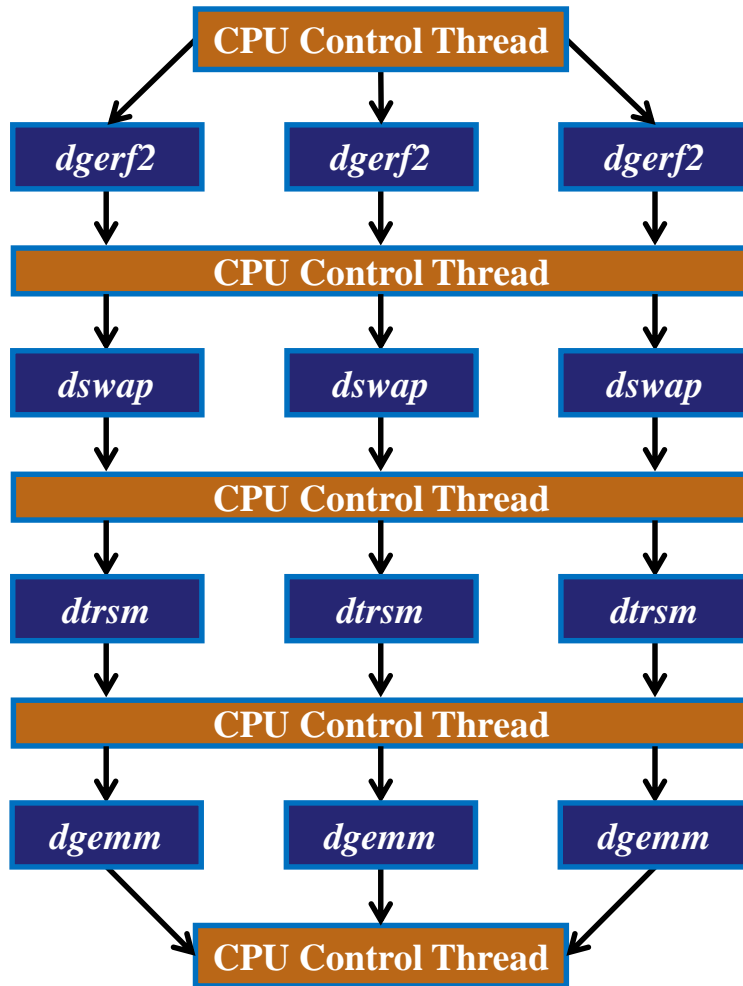
LU decomposition (Kepler)



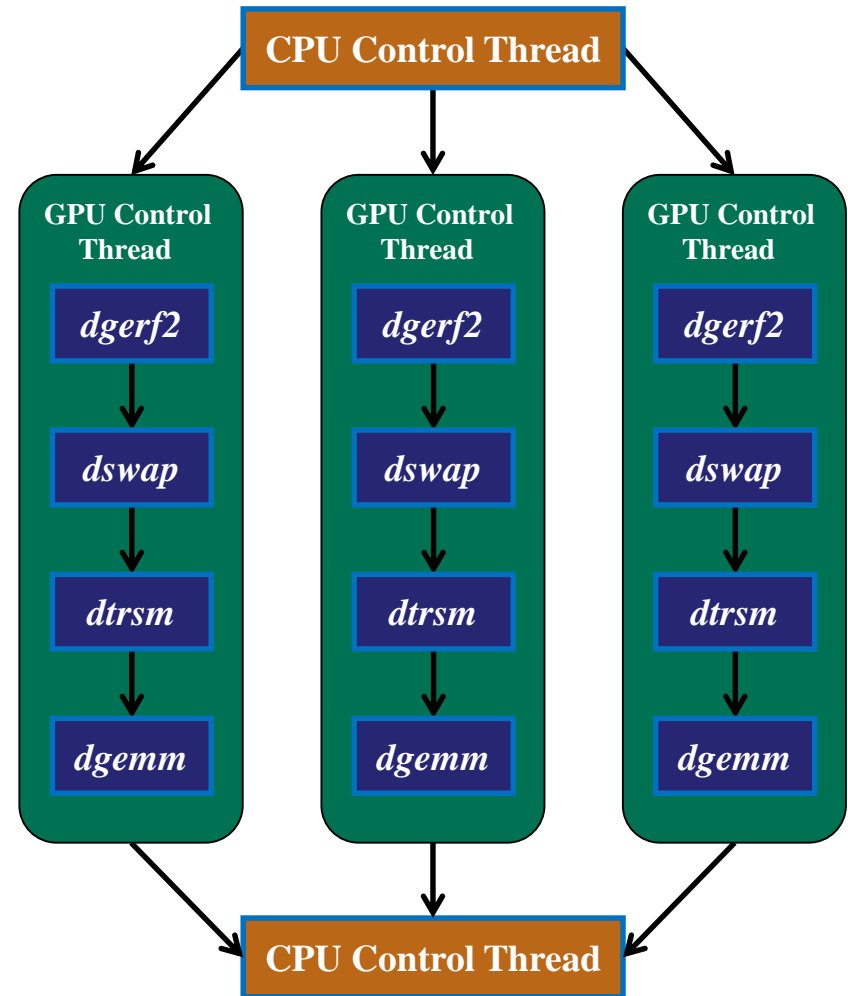
CPU Code

GPU Code

LU Decomposition



Multiple LU-Decomposition, Pre-Kepler



Batched LU-Decomposition, Pre-Kepler

4. Library Call

```
__global__ void libraryCall(float *a, float *b, float *c) {  
    // 所有线程都生成数据 All threads generate data  
    createData(a, b);  
    __syncthreads();  
    // Only one thread calls library  
    if(threadIdx.x == 0) {  
        cublasDgemm(a, b, c);  
        cudaDeviceSynchronize();  
    }  
    // 所有线程都等待 dt_rsm  
    __syncthreads();  
    // 继续  
    consumeData(c);  
}
```

cuBLAS:
Cuda矩阵运算库



5. Parallel Recursion

Quick sort

```
__global__ void qsort(int *data, int left, int right)
{
    int pivot = data[0];
    int *lptr = data + left, *rptr = data + right;
    // 根据中枢值将数据一分为二
    partition(data, left, right, lptr, rptr, pivot);
    // 启动下一级递归
    if (left < (rptr - data))
        qsort<<<...>>>(data, left, rptr - data);
    if (right > (lptr - data))
        qsort<<<...>>>(data, lptr - data, right);
}
```

